



# **T03: BVM2025 Hands-on Tutorial on Implicit Neural Representations (INRs) in Medical Imaging**

Mattias, Fenja, Ziad and Christoph | University of Lübeck

Sunday 09 March 2025, 14:00 am - 17:20 pm (CET)

<https://ziadhemidi.github.io/INR4BVM25.github.io/>

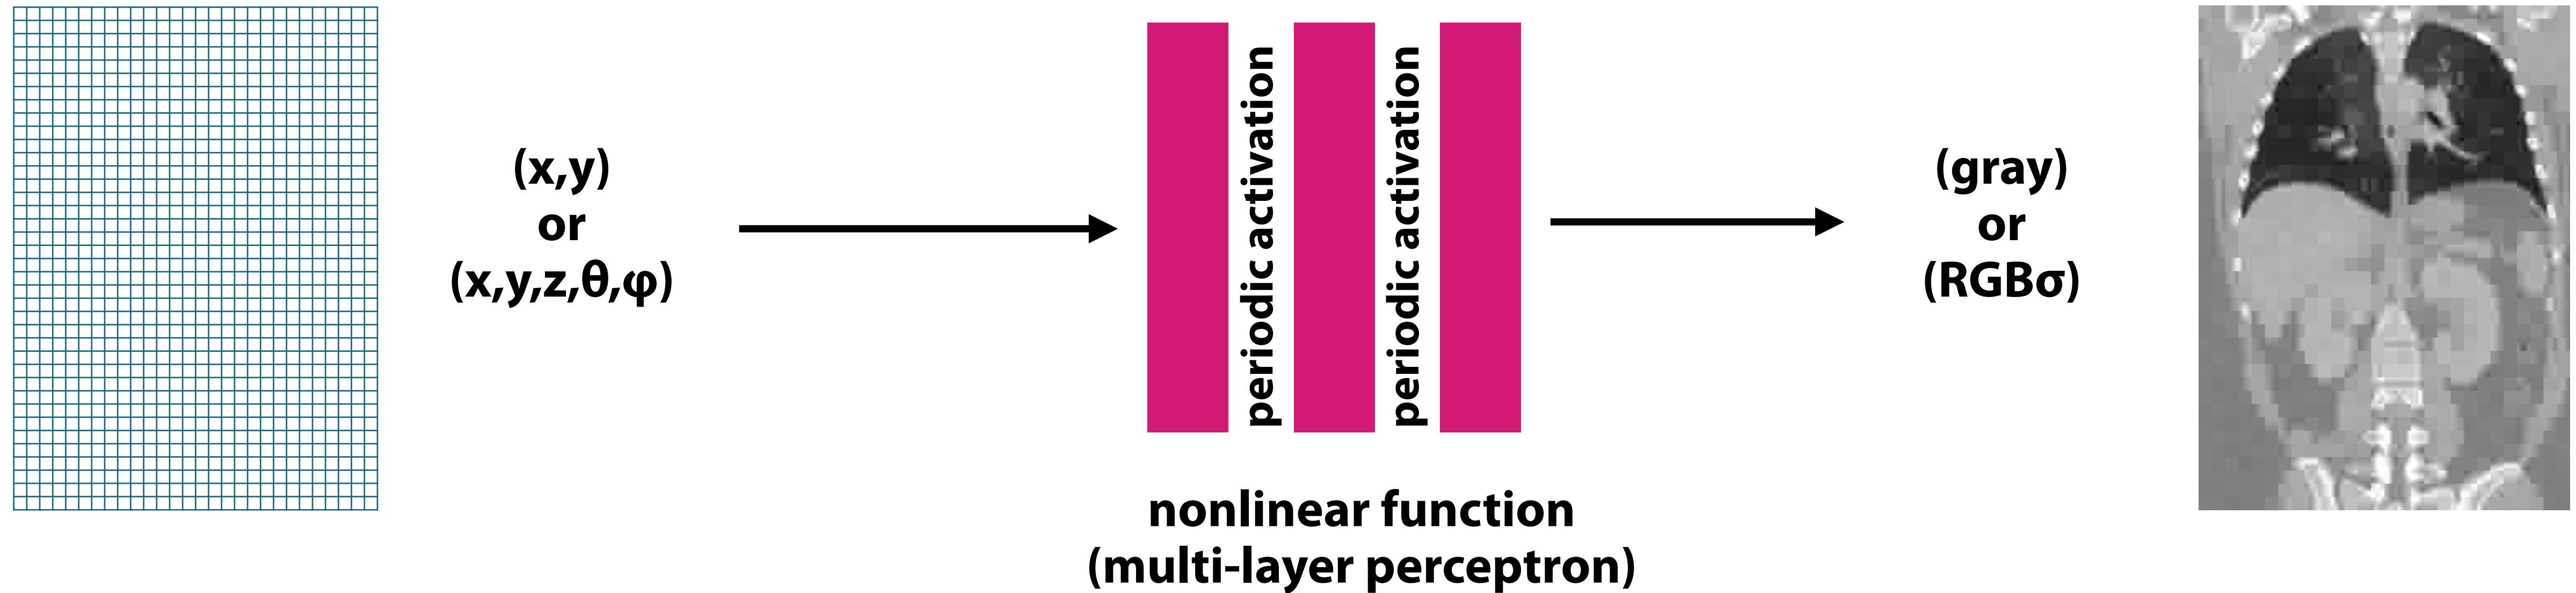


# Schedule for today

Time	Topic	Contents
14:00 - 14:10	<b>Welcome &amp; Overview</b>	Introduction of speakers and organisational tutorial information
14:10 - 14:30	<b>Introduction to INR</b>	Basic conceptual and mathematical foundations on INRs and associated activation functions
14:30 - 16:30	<b>Hands-On (incl. individual breaks)</b>	Assisted programming of INR solutions for reconstruction, denoising and image registration in Jupyter notebooks
16:30 - 17:10	<b>Advanced Topic Mini-Lectures (4x10min)</b>	10 Minute Mini-Lectures on challenging MRI reconstruction, rendering INRs (NERF), implicit neural image registration and multi-instance INR generalization
17:10 - 17:30	<b>Evaluation</b>	Discussion / Q&A

# Key idea: encode image as nonlinear function (MLP)

**often also called coordinate MLP**



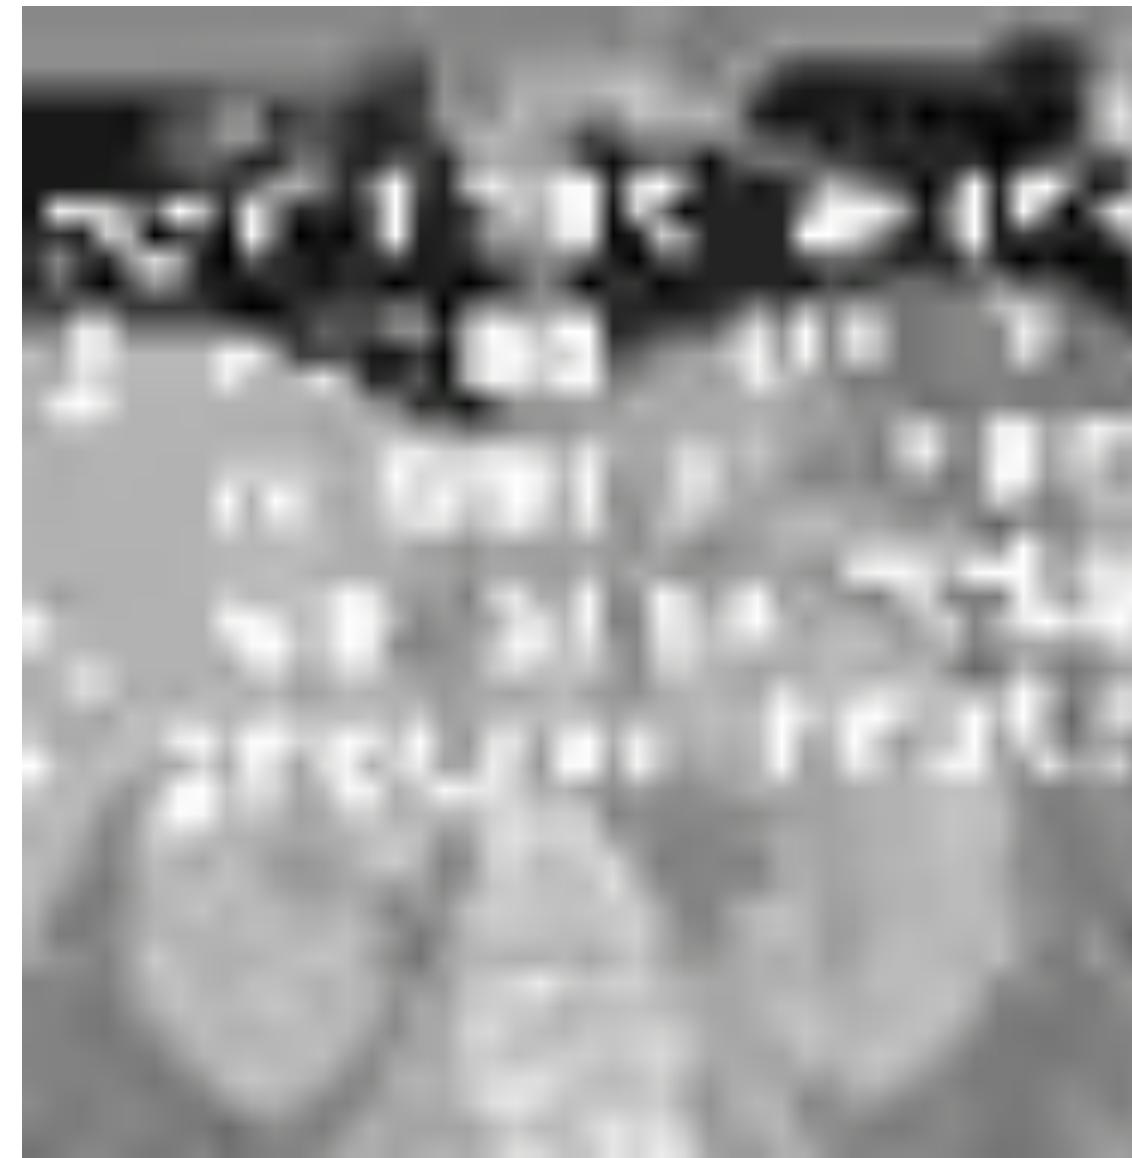
- **input** to neural field is a set of  **$x,y$  coordinates** (+ potential Fourier encoding)
- **output are gray/colour values** or additionally depth/alpha-transparency
- **multi-layer perceptron** is used to **learn functional mapping** from xy coordinates to gray-value

**"training"** on single input

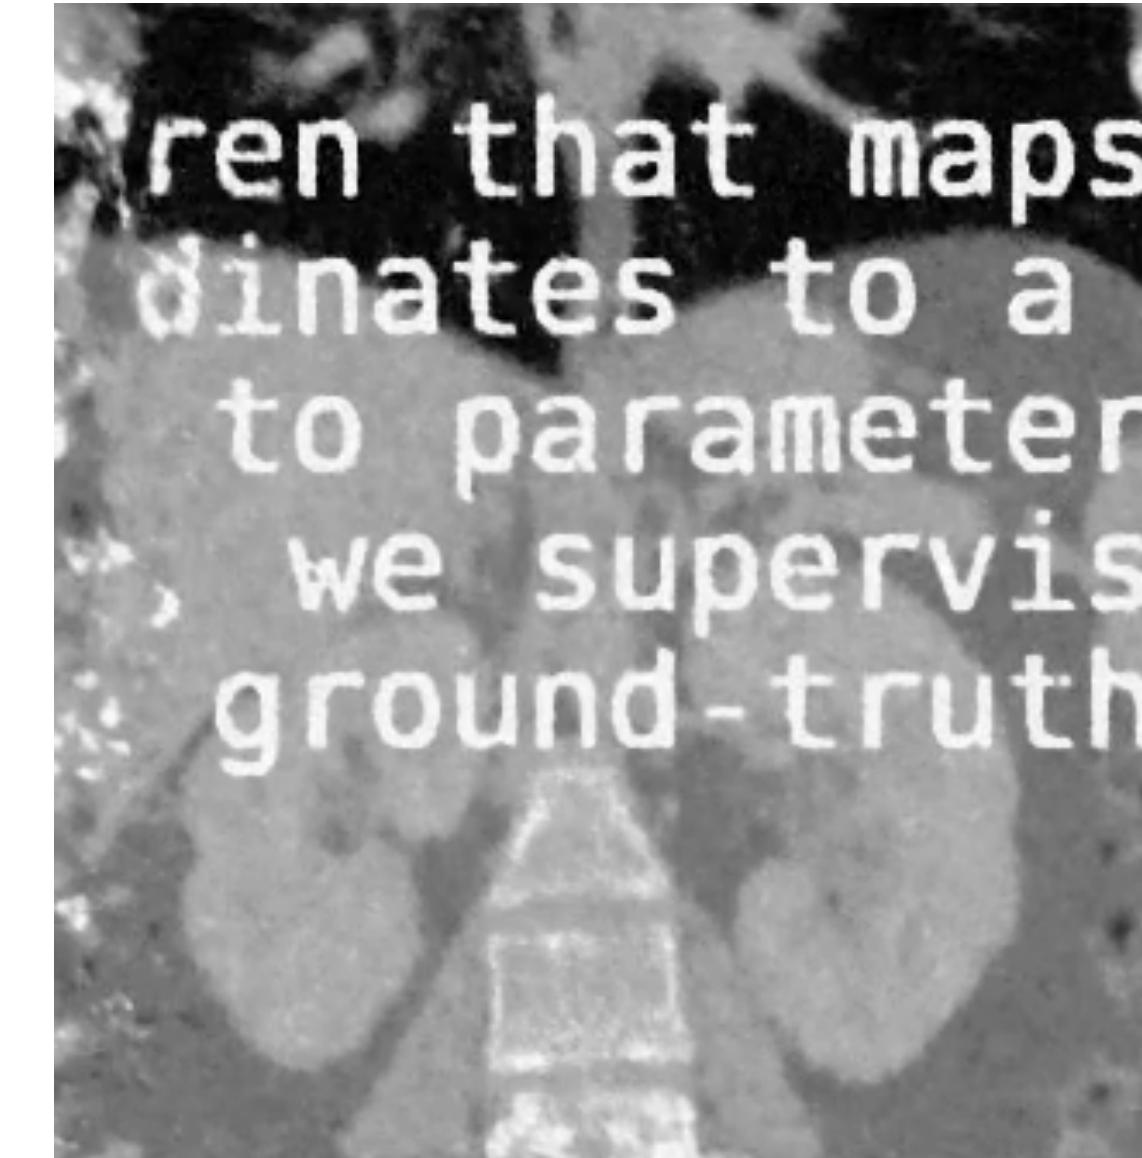
can be used **out-of-the-box**  
**optimisation** for superresolution /  
denoising / restoration

**periodic activation is**  
**important**

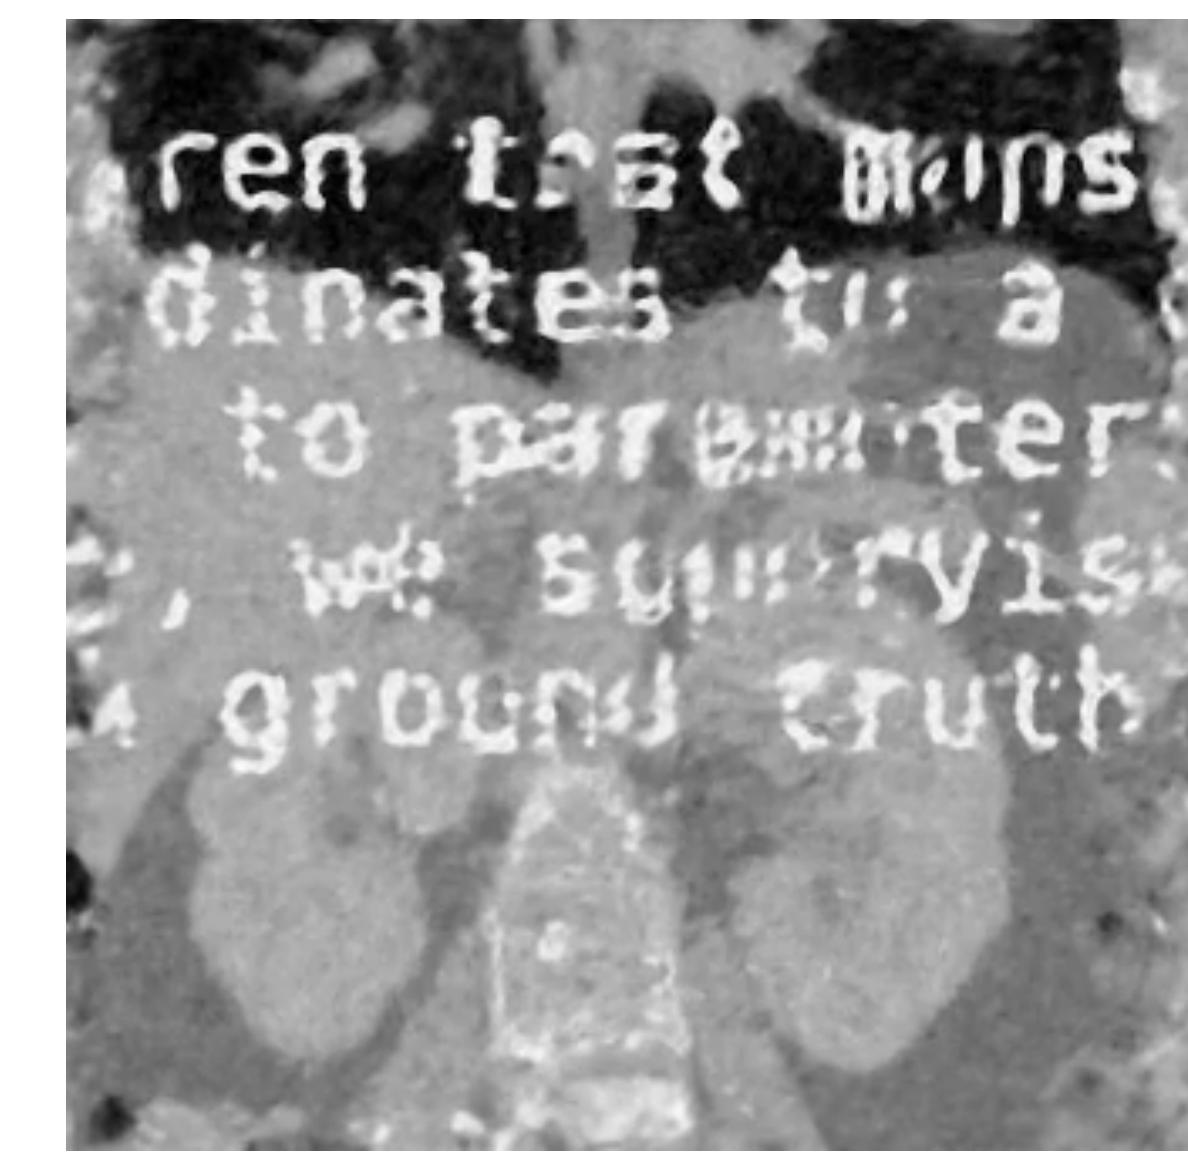
# Motivation/Teaser I: multi-view superresolution



**input sequence is low resolution (84x84 pixels)**  
"camera position" of 64 frames is known

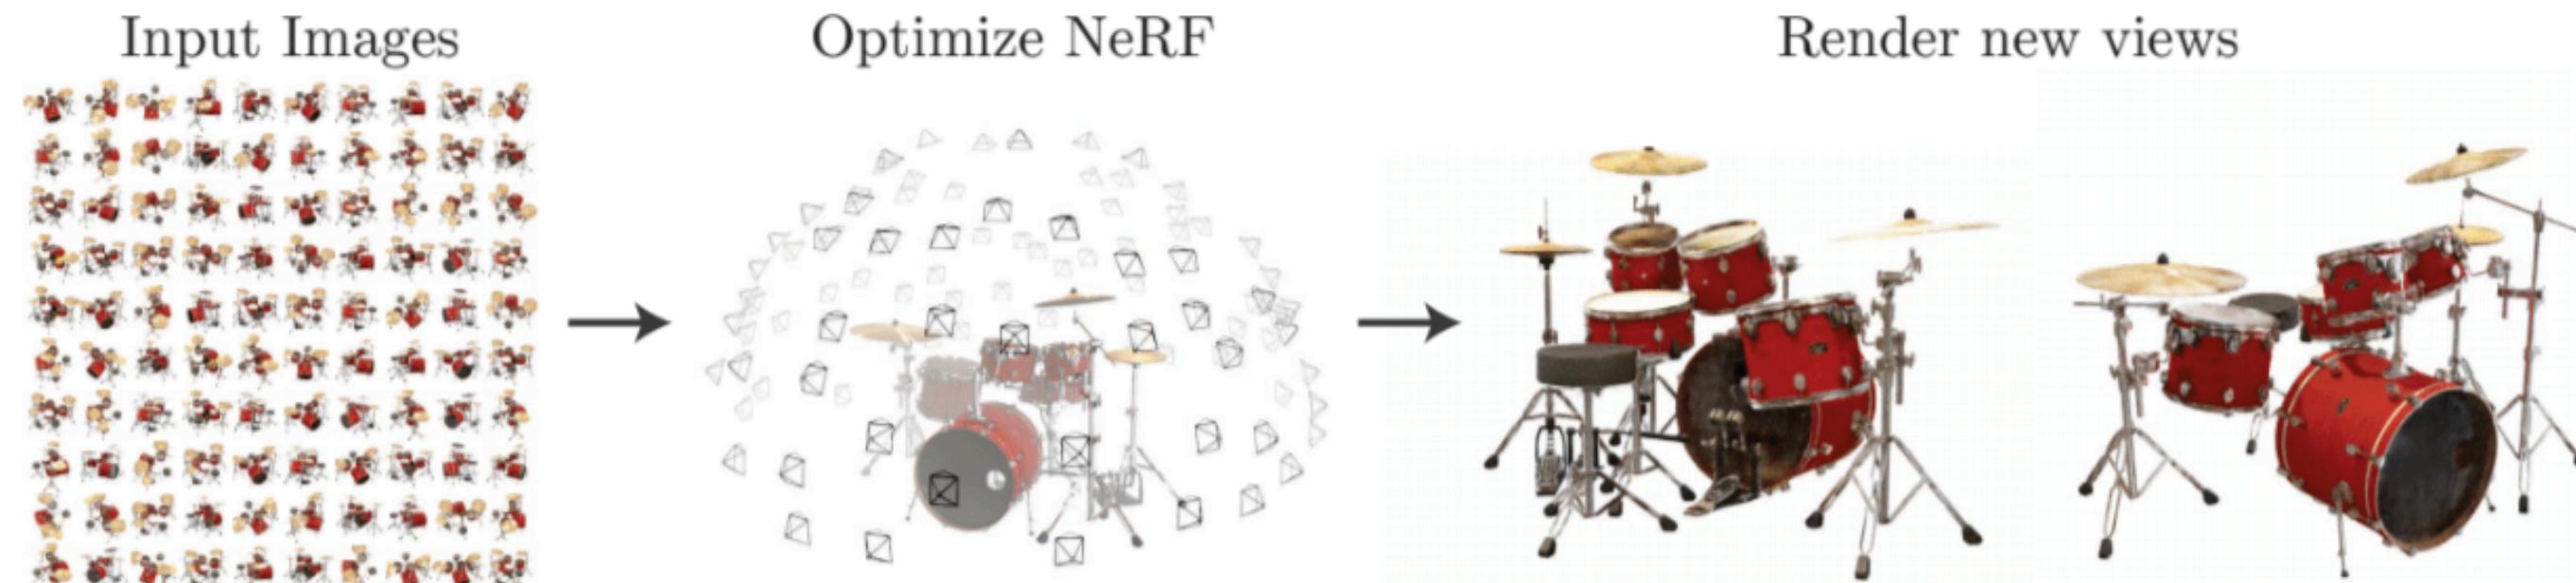


SIREN coordinate-MLP  
**reconstructs 8x super-resolved image**  
with dimensions 334x334



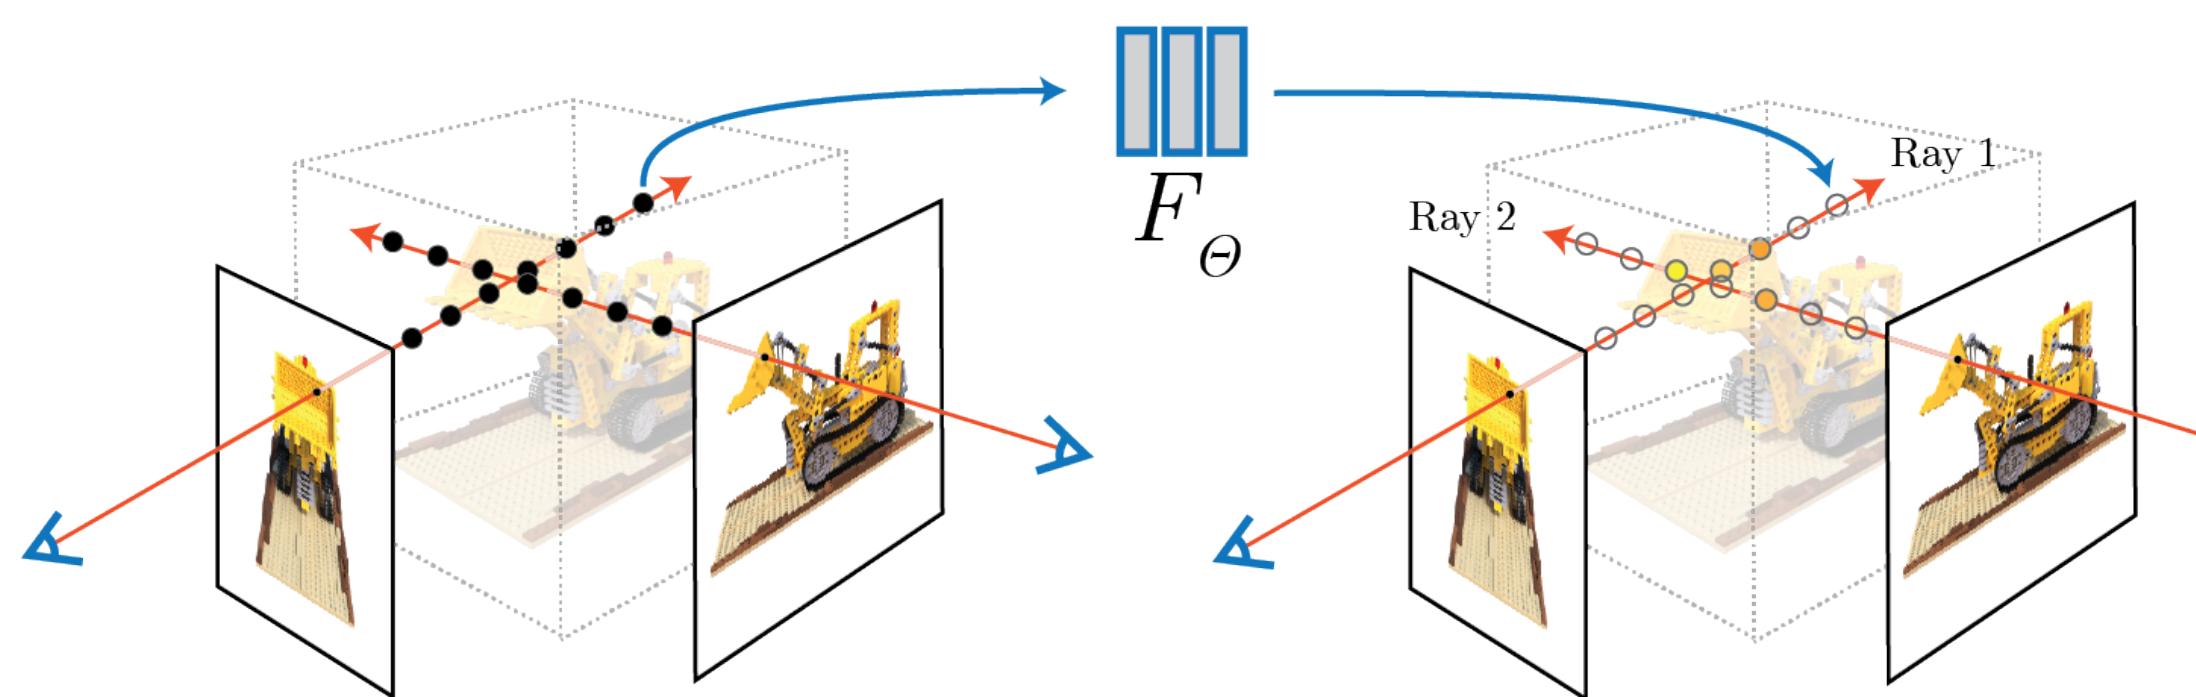
even first 8 frames of low-resolution sequence are sufficient to make most of the text readable

# Motivation/Teaser II: Neural Radiance Fields

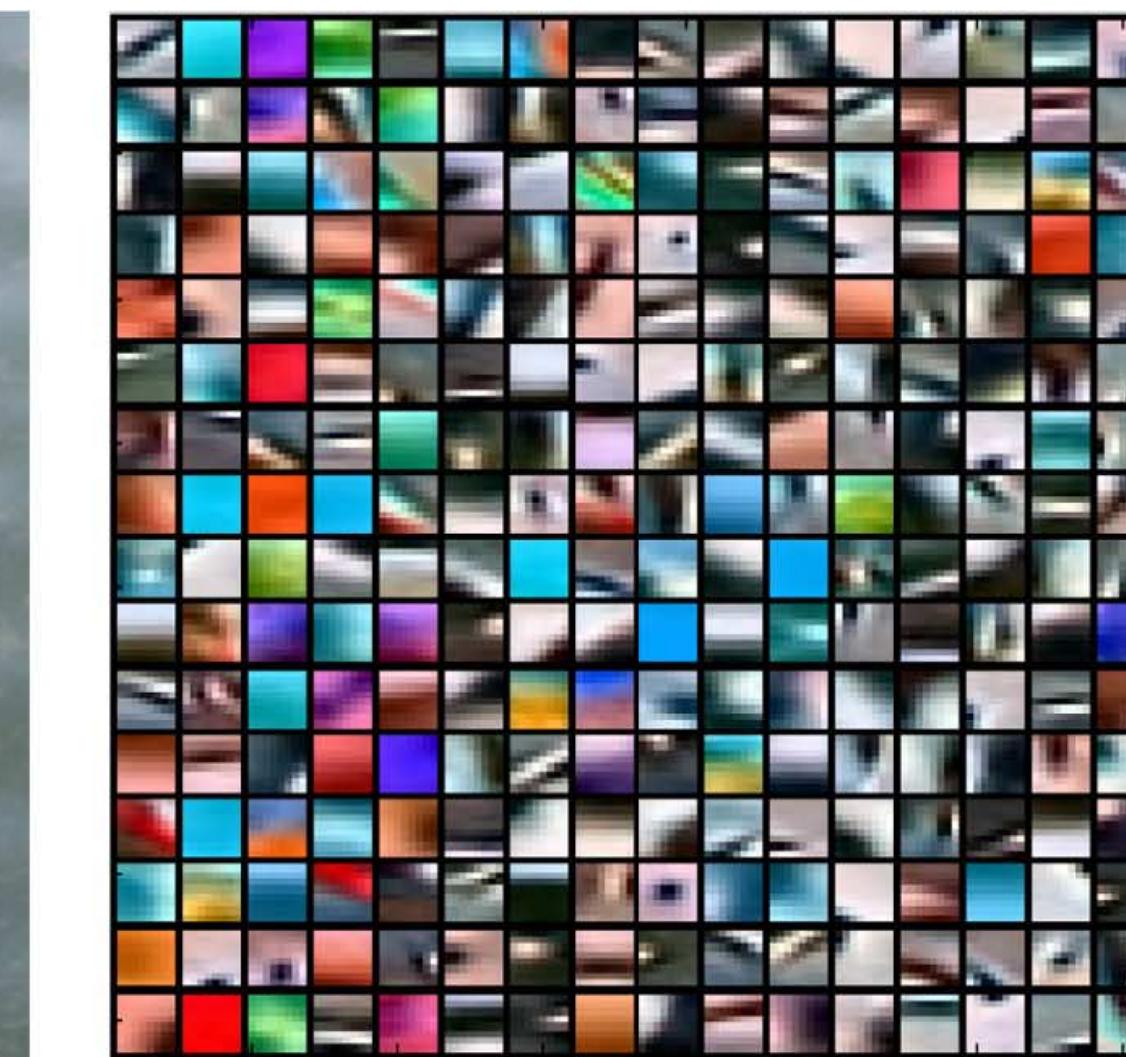
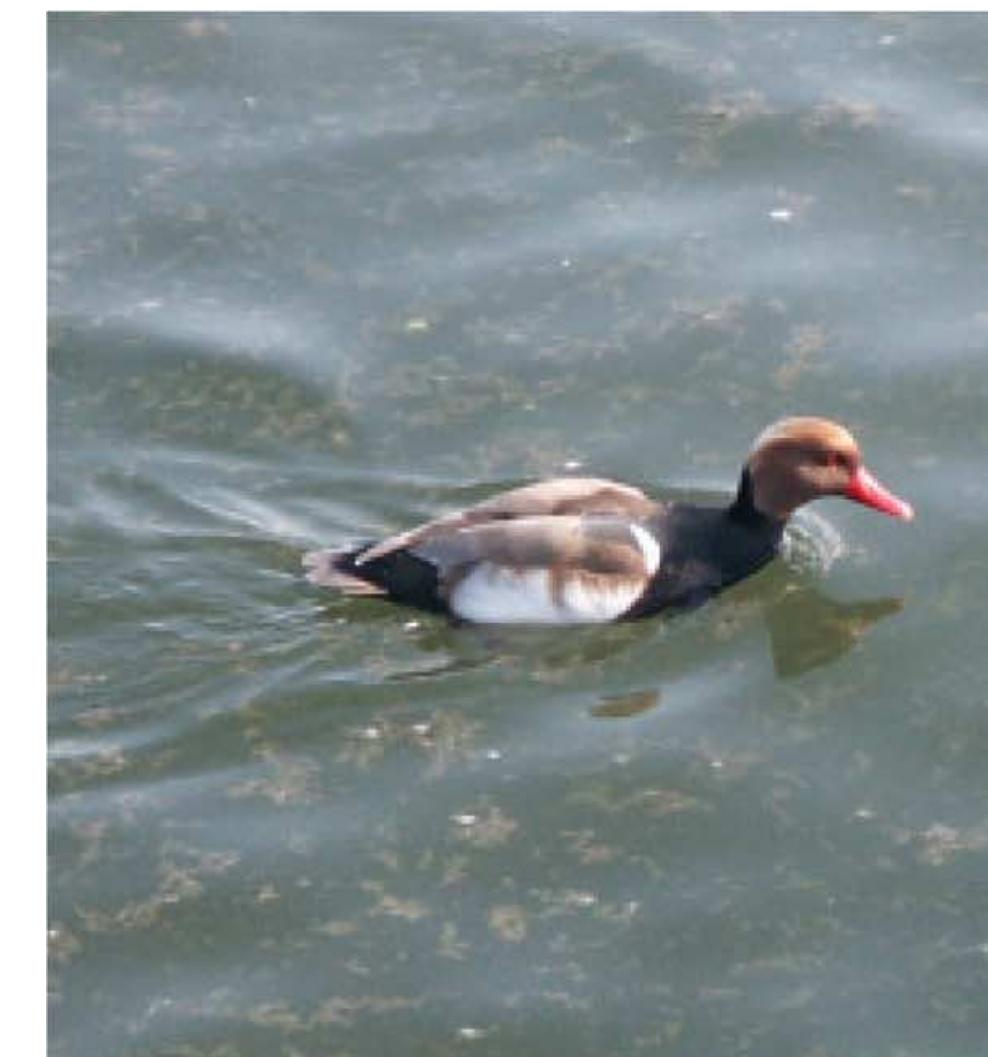
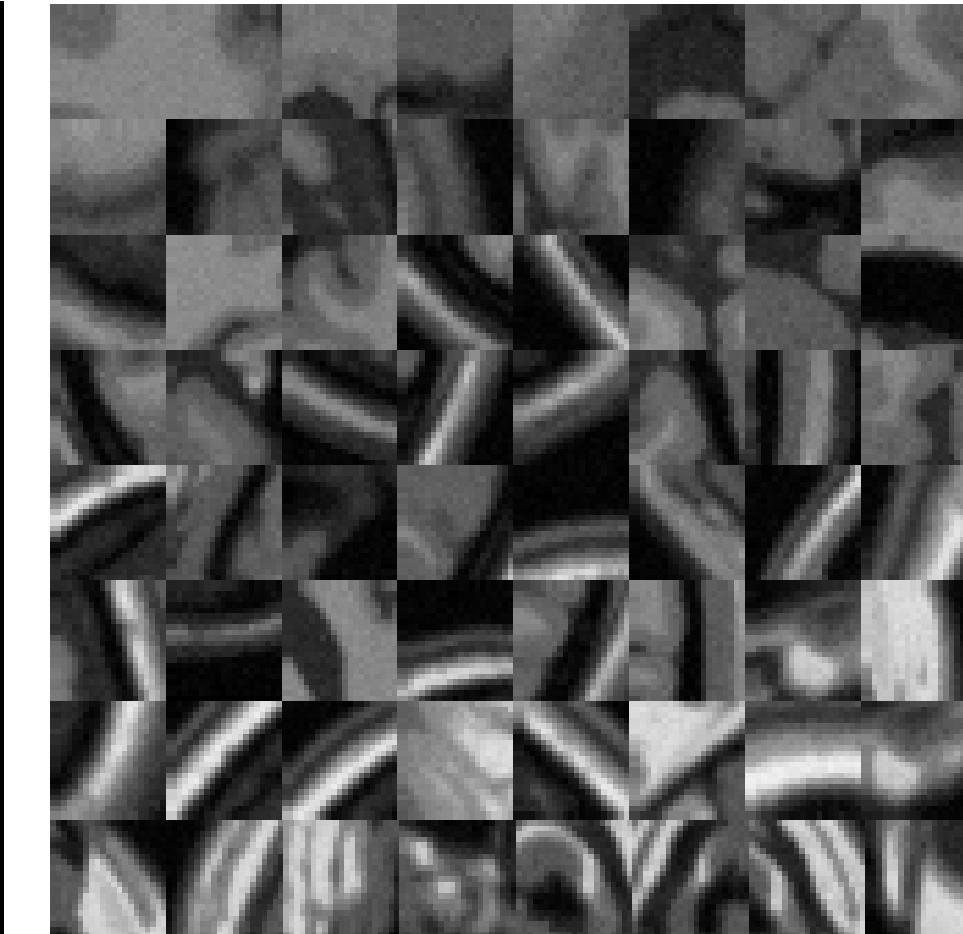
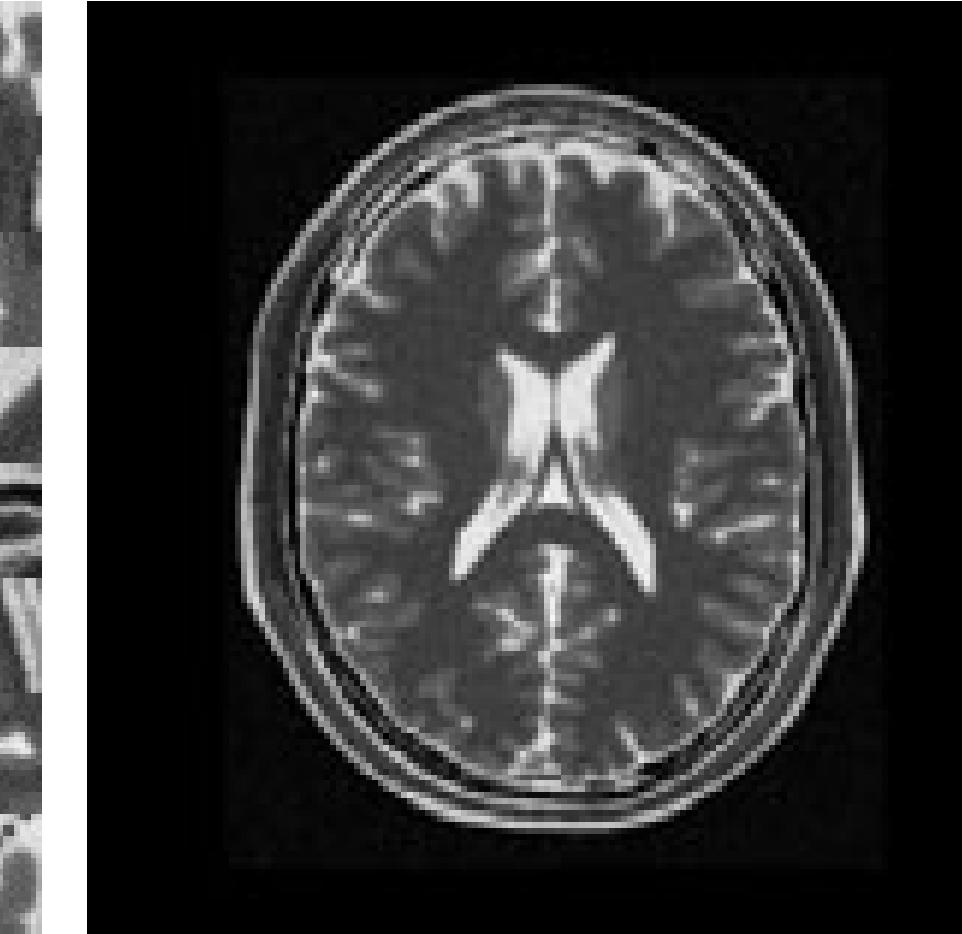
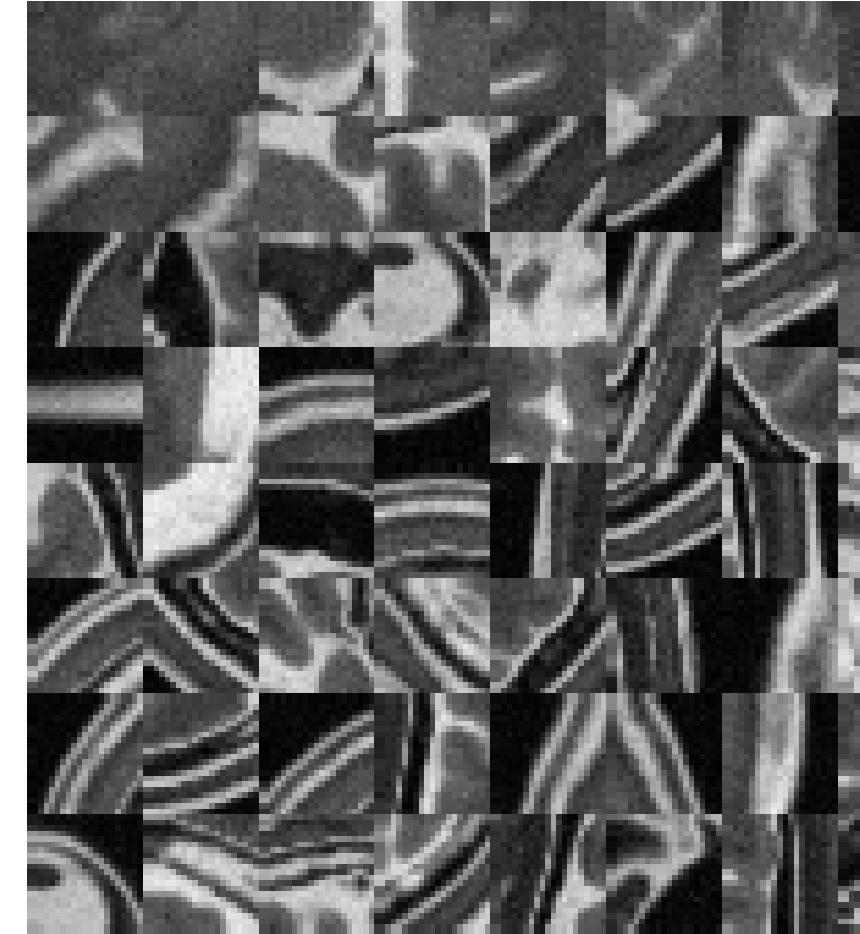


A NeRF stores a volumetric scene representation as the weights of an MLP, trained on many images with known pose.

**New views are rendered** by integrating the density and color at regular intervals along each viewing ray.



# Representative elements of images



Yi, Soatta: "Multimodal Registration via Spatial-Context Mutual Information" IPMI 2011  
Mairal et al.: "Sparse Representation for Color Image Restoration" IEEE TIP 2008



# Periodic Image/Coordinate Representations

# Discrete cosine transformation (DCT)

Reminder: Fourier transform

Represents (complex) image by frequency/phase spectrum

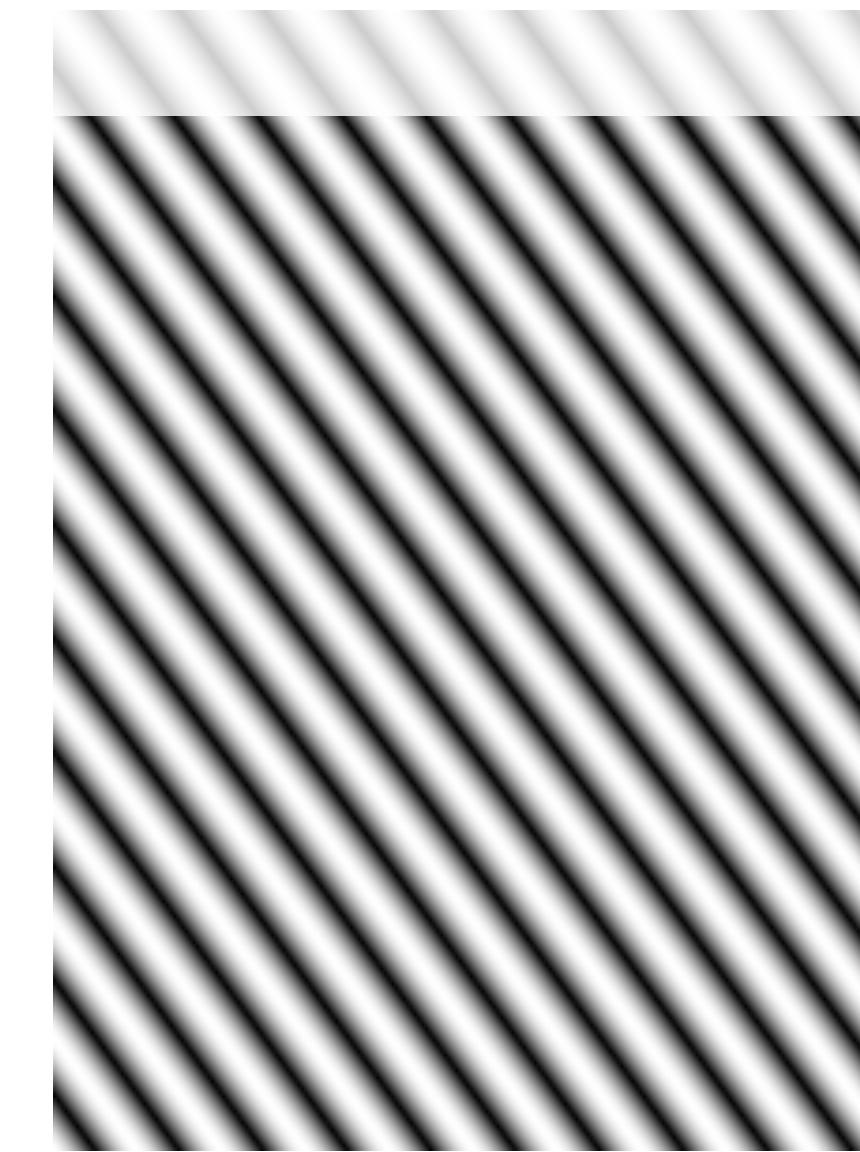
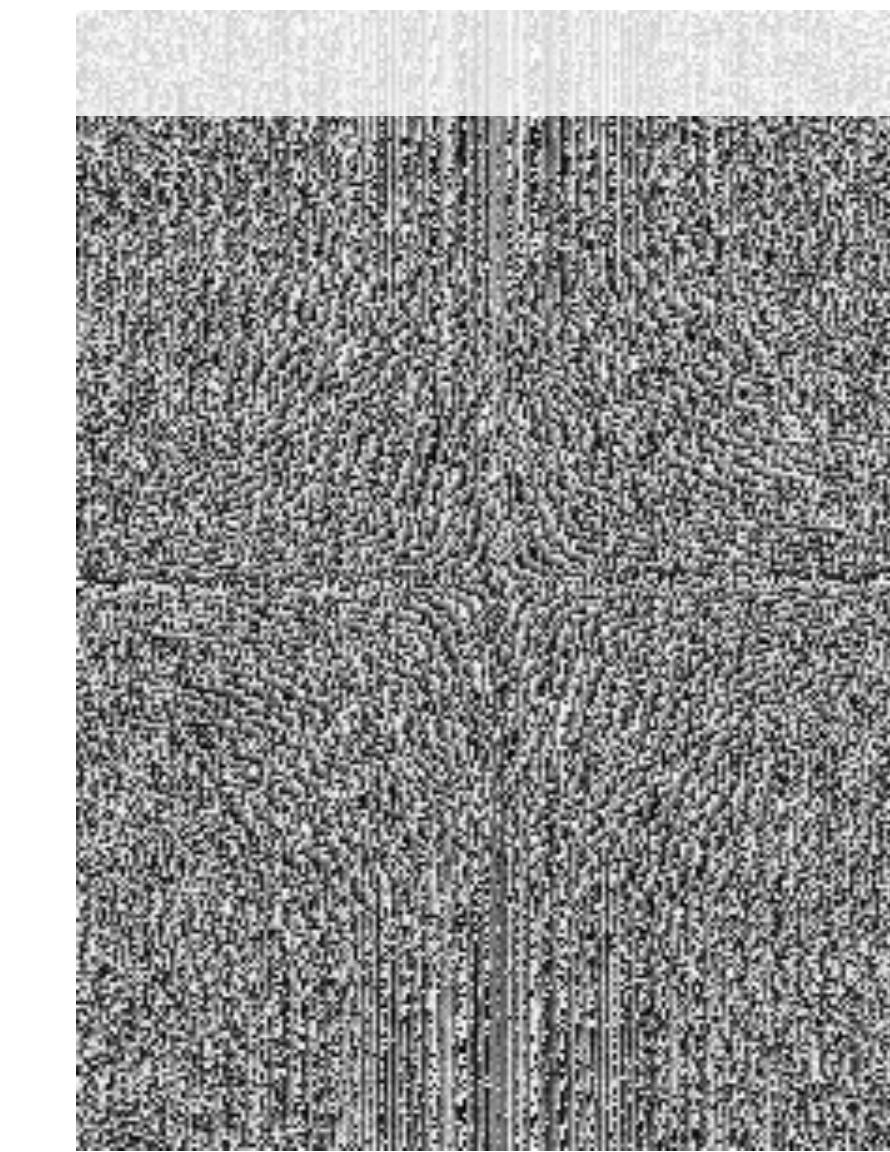
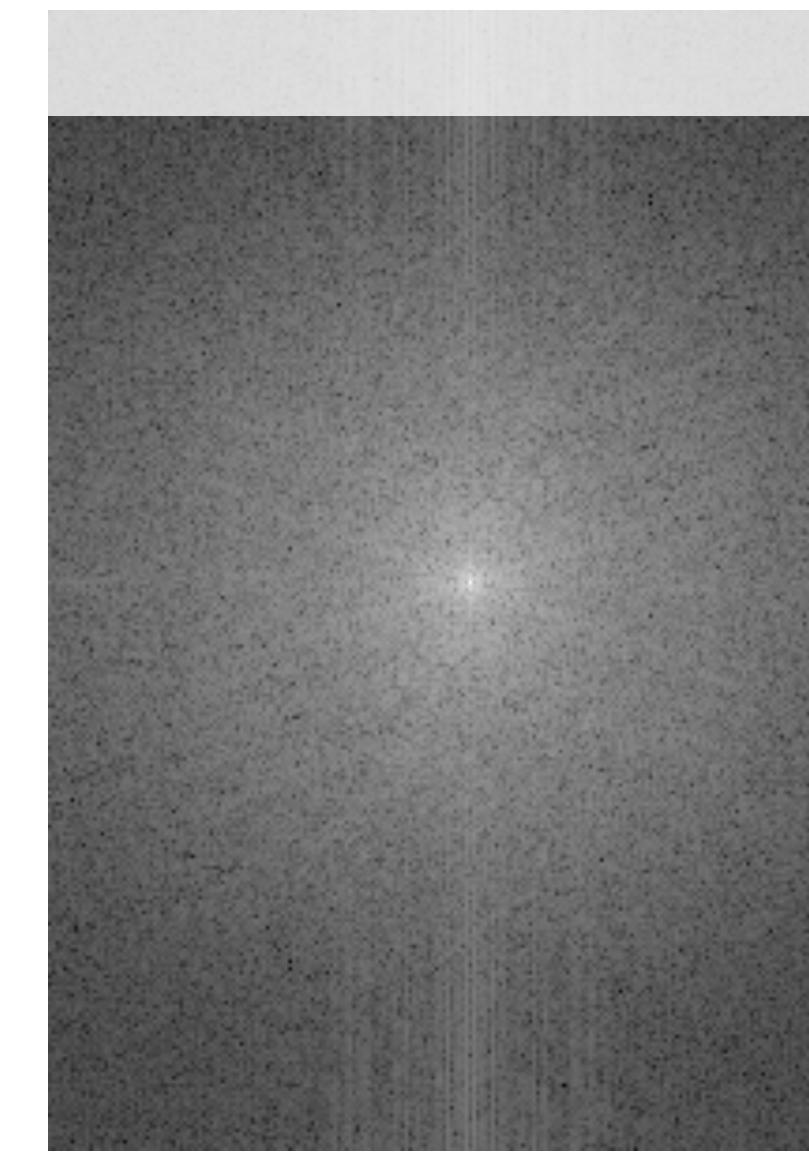
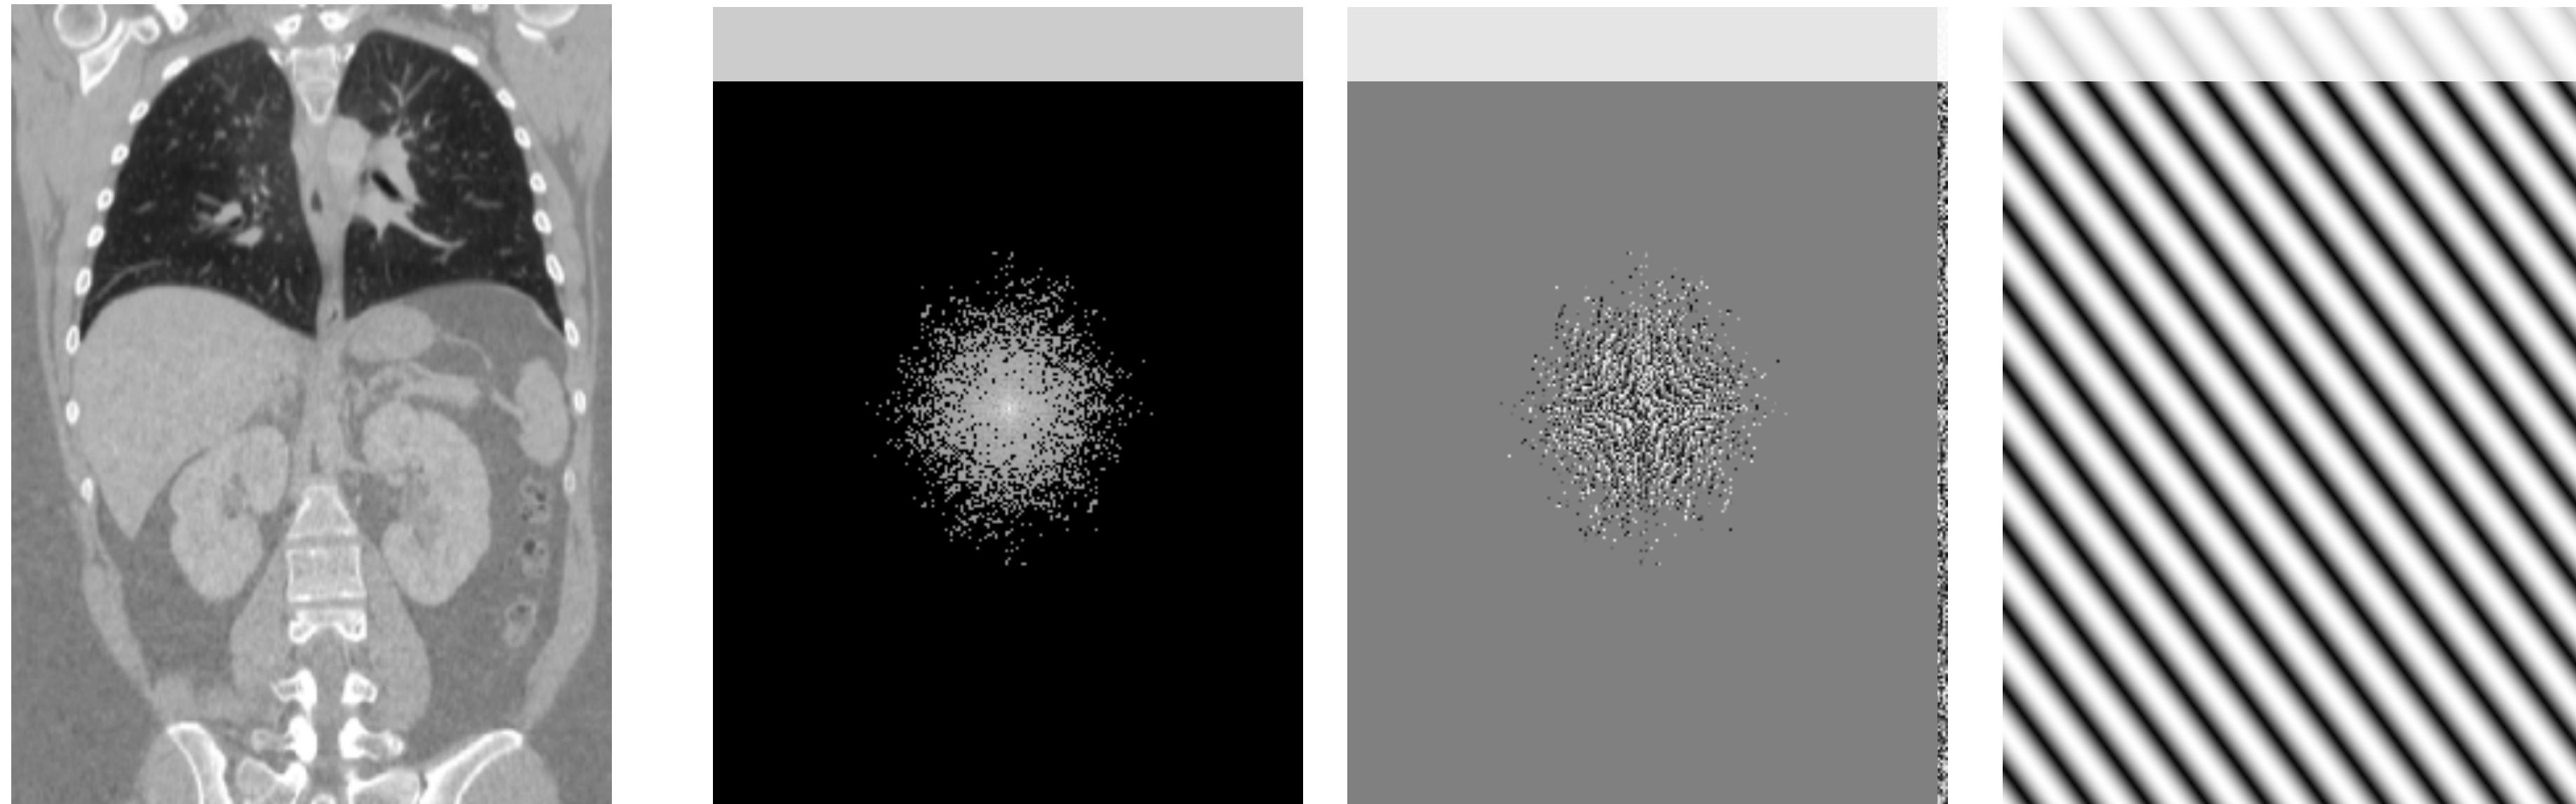


Image is created as a superposition of 2D sine waves  
Reconstruction of the original image possible without loss

# Discrete cosine transformation (DCT)

Reminder: Fourier transform

Represents (complex) image by frequency/phase spectrum



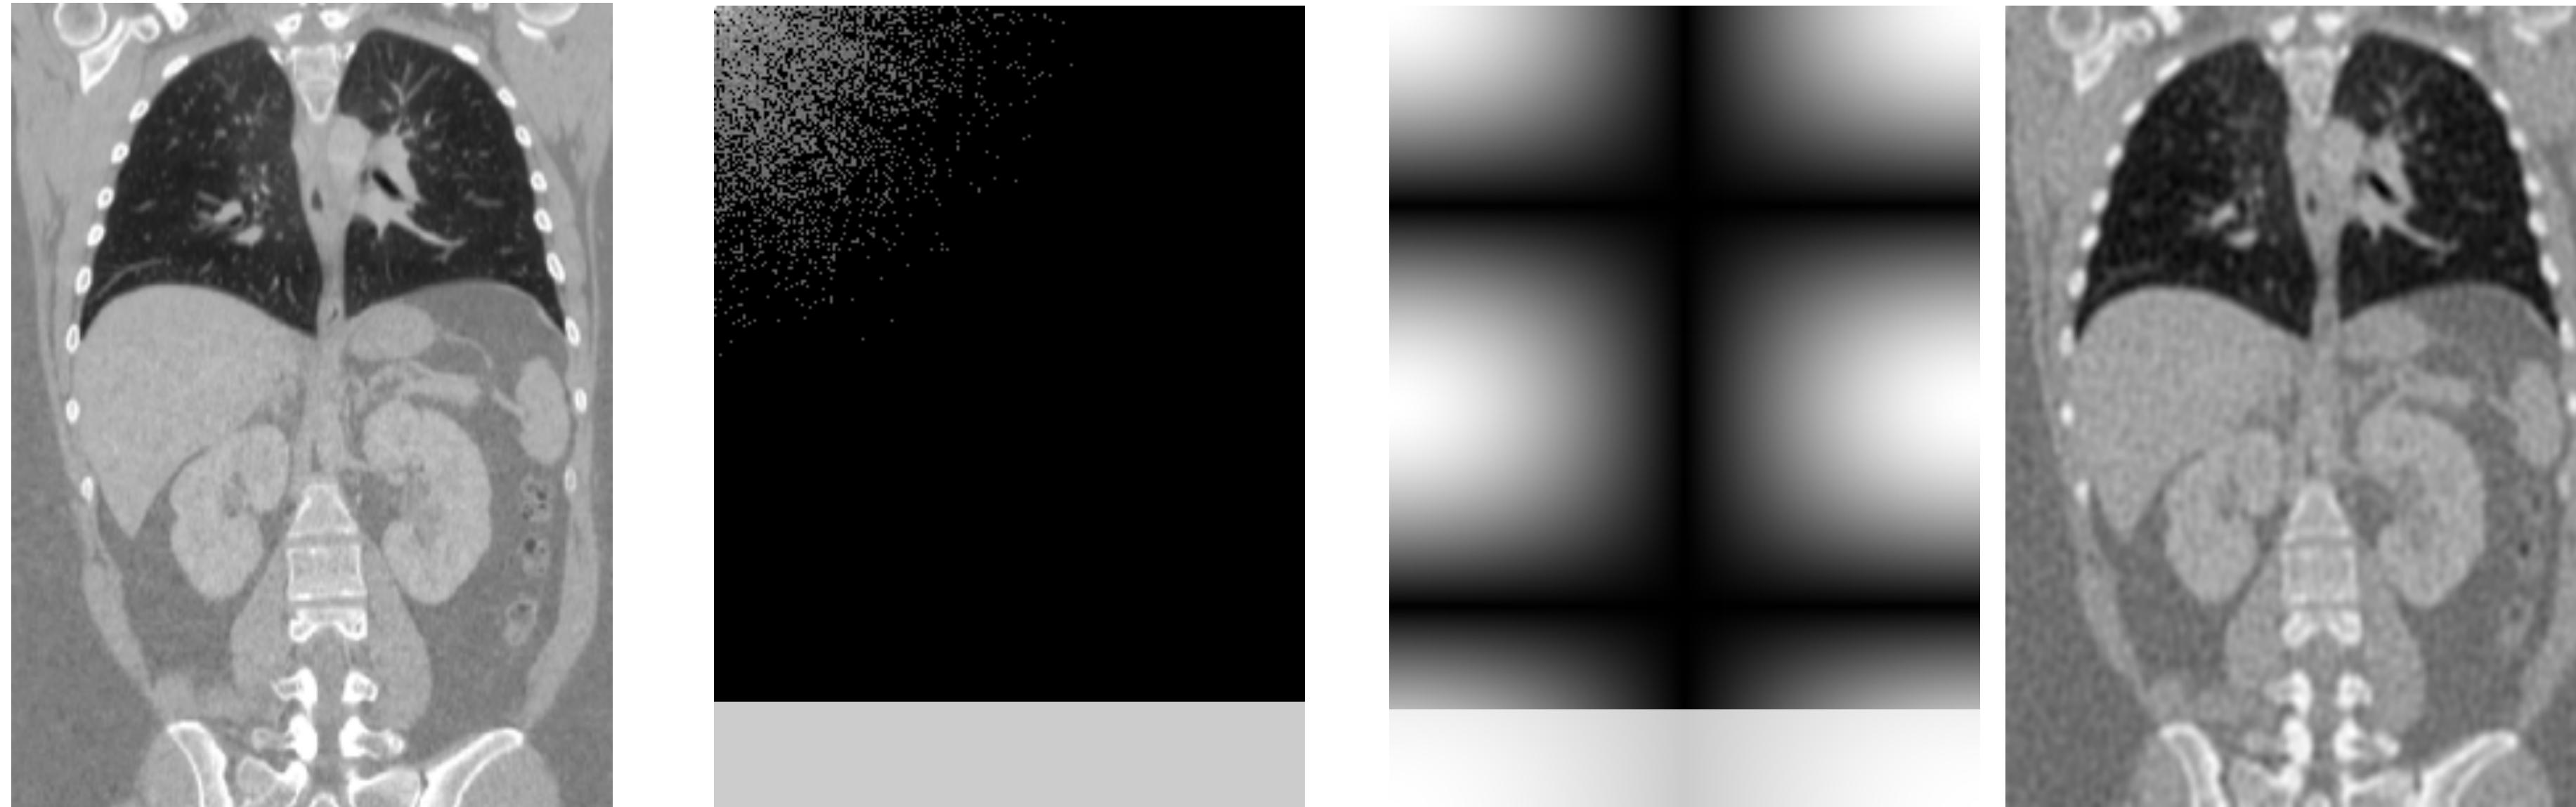
Lossy compression in Fourier space

Discard ( $=0$ ) of the 95% smallest coefficients in terms of magnitude

# Discrete Cosine Transformation (DCT)

DCT does not contain an imaginary part

Forward transformation 1D: 
$$X_k = \sum_{n=0}^{N-1} x_n \cos \left[ \frac{\pi}{N} \left( n + \frac{1}{2} \right) k \right] \quad k = 0, \dots, N - 1.$$



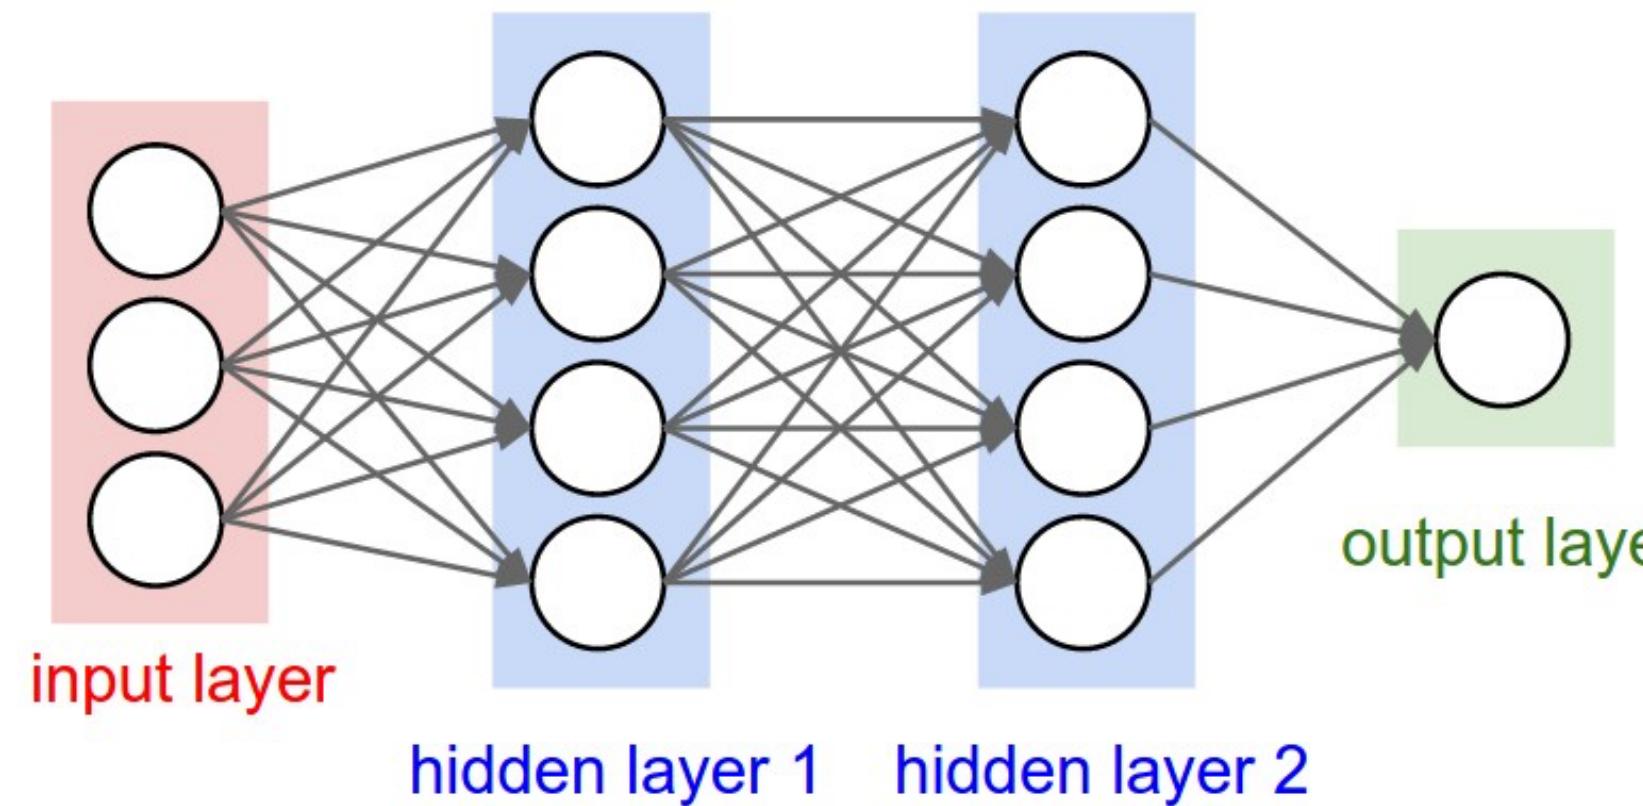
Transformation with known dimensions can be represented as a twofold matrix product



# Representing images as learned functions

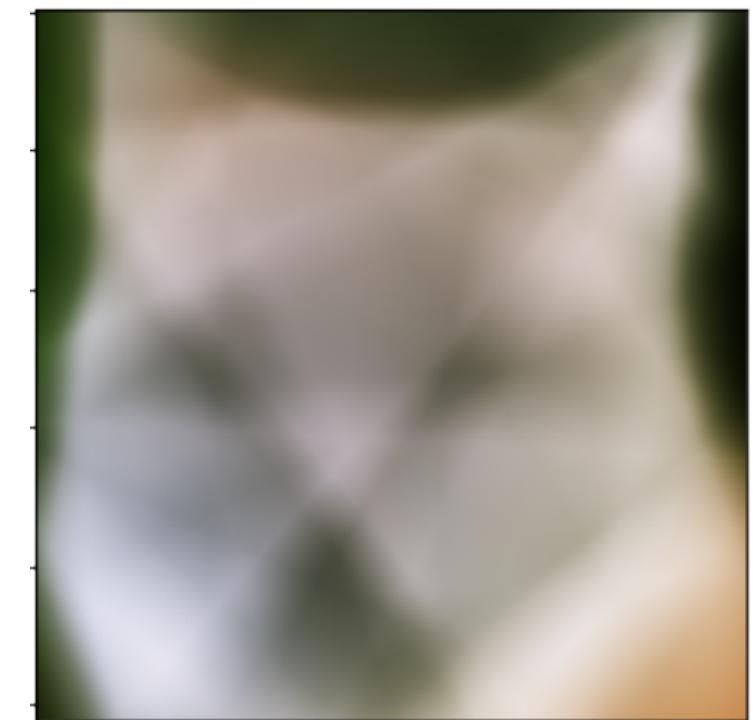
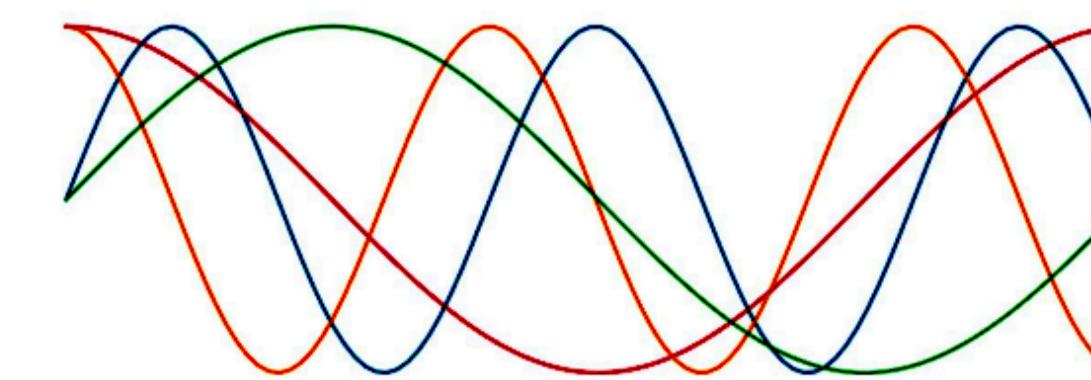
# Representing images as **learned** functions

- **what if we could use non-linear functions?**



it is **important to mitigate low-frequency bias of ReLU MLP**

solution I: Coordinate Encoding / Projection

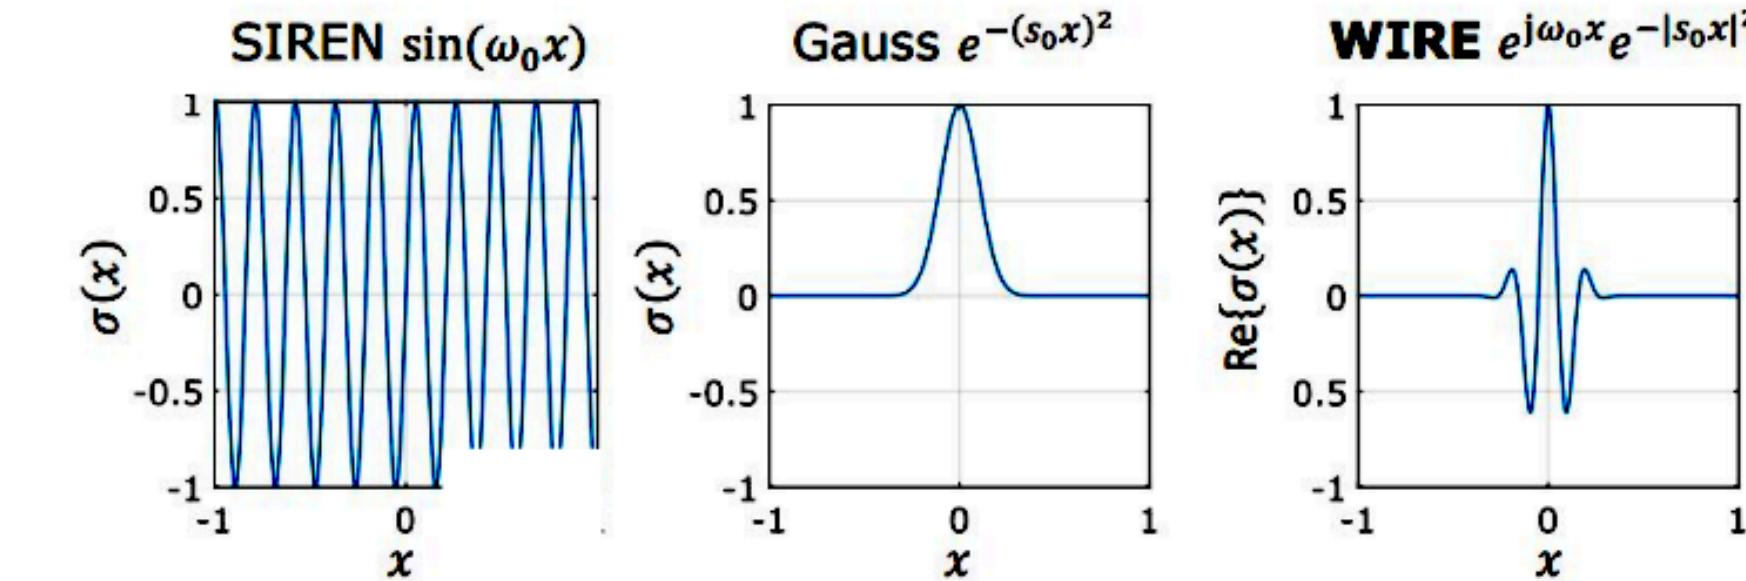


**basic coordinate  
ReLU-MLP**

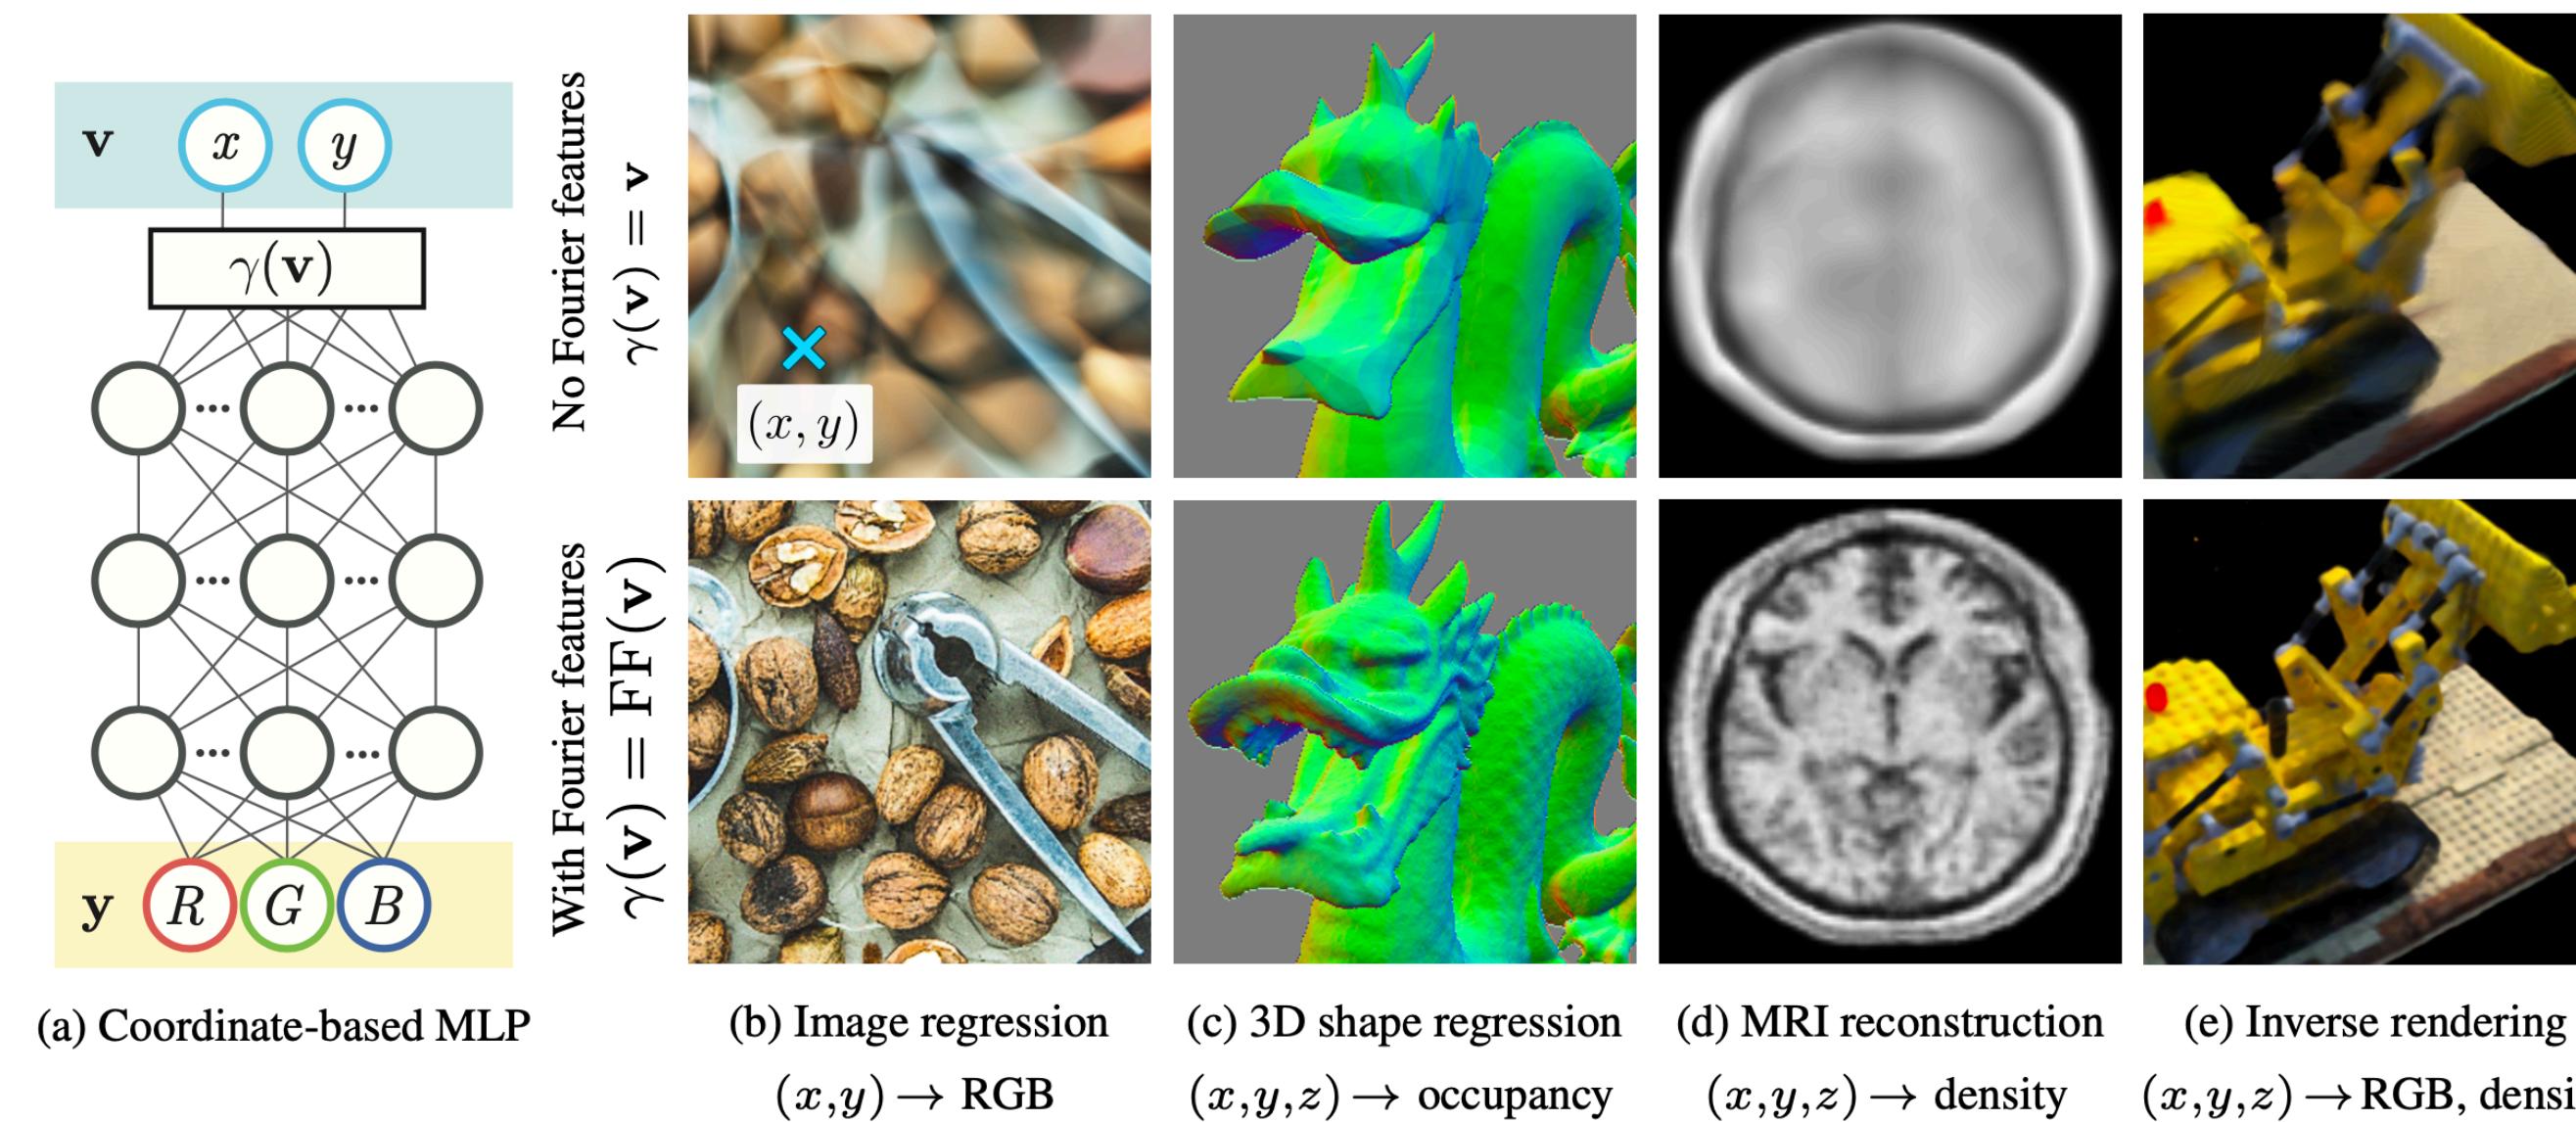
why did no one use this before 2020? Problem →  
**MLPs** with ReLU or tanh/sigmoid as activations **are surprising bad at learning periodic functions**

- **solution I:** use Fourier features (similar to cosine transform)
- **solution II:** use sinusoidal or exponential activation

solution II: Non-ReLU Activation Functions



# Representing images as functions of Fourier features



(a) Coordinate-based MLP

(b) Image regression

 $(x,y) \rightarrow \text{RGB}$ 

(c) 3D shape regression

 $(x,y,z) \rightarrow \text{occupancy}$ 

(d) MRI reconstruction

 $(x,y,z) \rightarrow \text{density}$ 

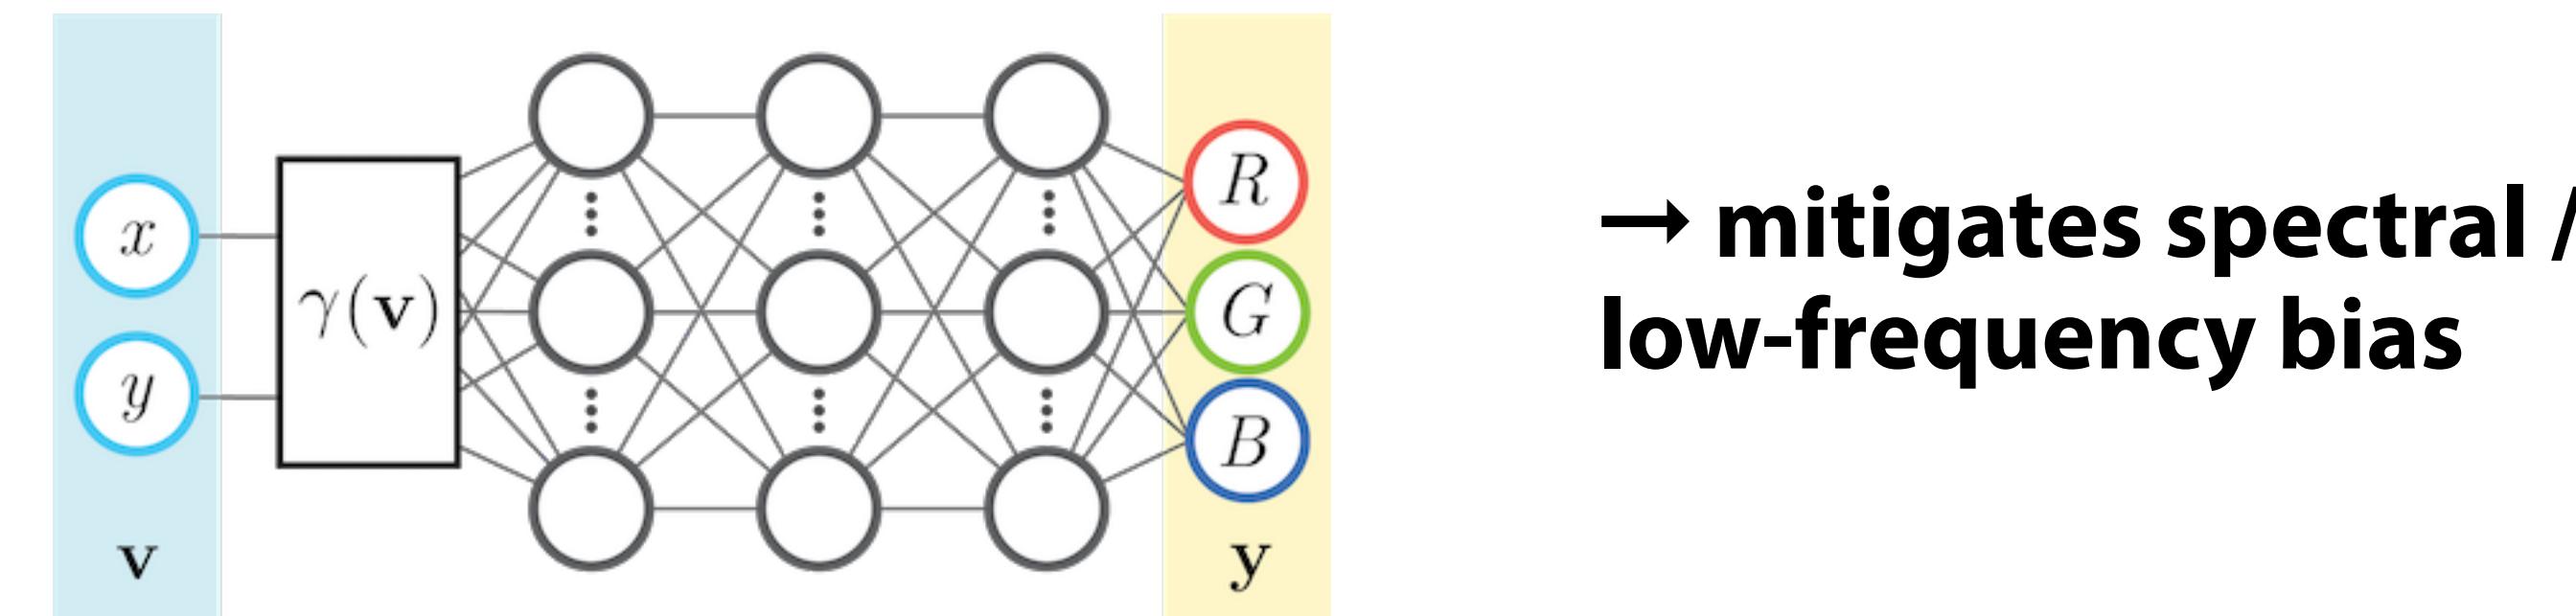
(e) Inverse rendering

 $(x,y,z) \rightarrow \text{RGB, density}$ 

Matthew Tancik et al.: **Fourier Features Let Networks Learn High Frequency Functions in Low Dimensional Domains**  
NeurIPS 2020 (spotlight)

# Fourier feature mapping

- as mentioned in teasers: **the idea is to represent the mapping from xy-coordinates to RGB-values as multi-layer perceptron (MLP)**



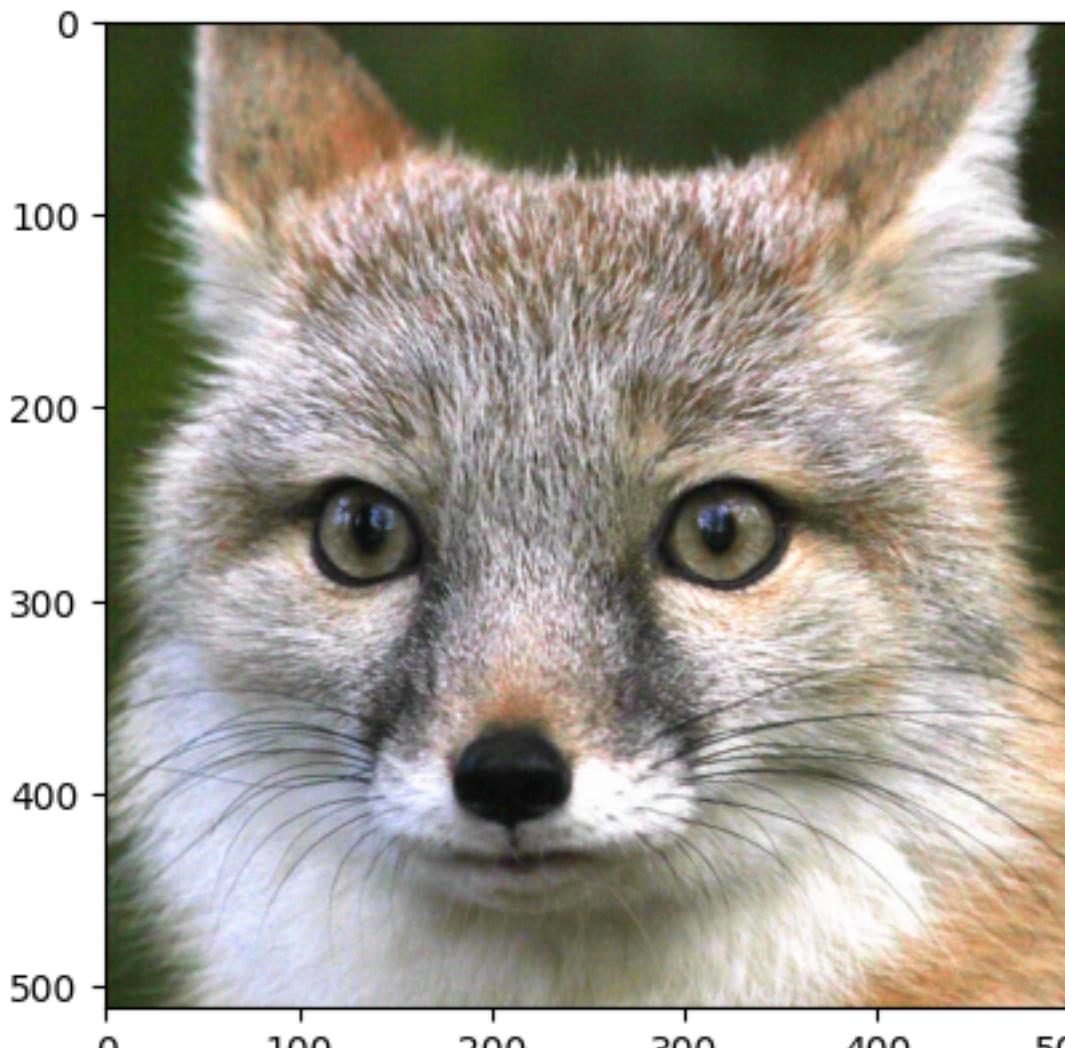
**(Gaussian) Fourier feature mapping → transforms input coordinates**

$$\gamma(\mathbf{v}) = [\cos(2\pi \mathbf{B}\mathbf{v}), \sin(2\pi \mathbf{B}\mathbf{v})]^T$$

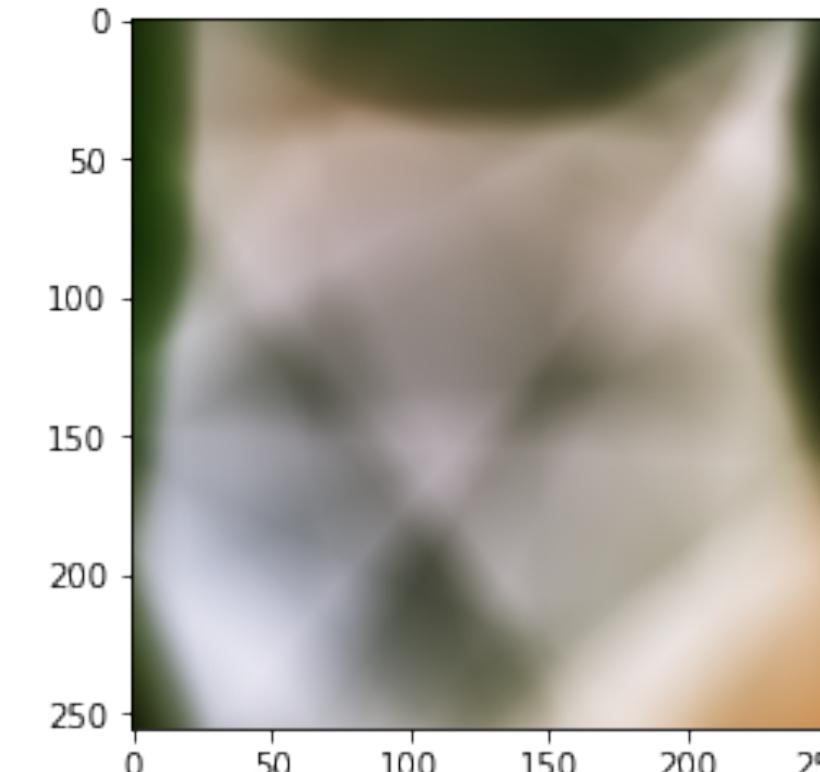
each entry in  $\mathbf{B} \in \mathbb{R}^{m \times d}$  is sampled from  $\mathcal{N}(0, \sigma^2)$

**d is usually 32-64**

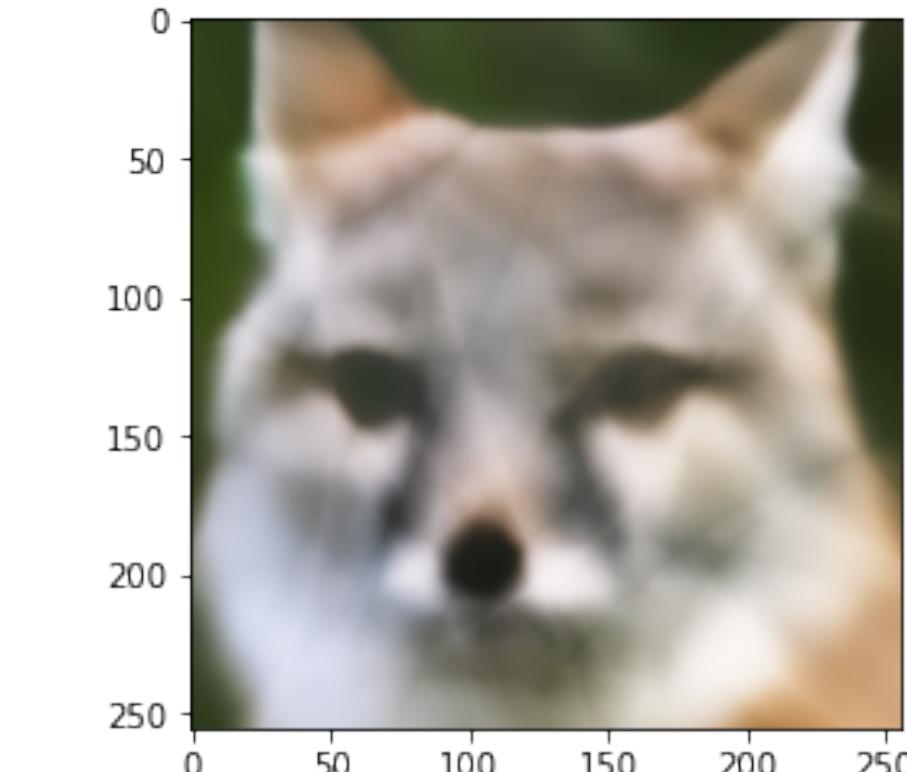
# Representing images as functions of Fourier features



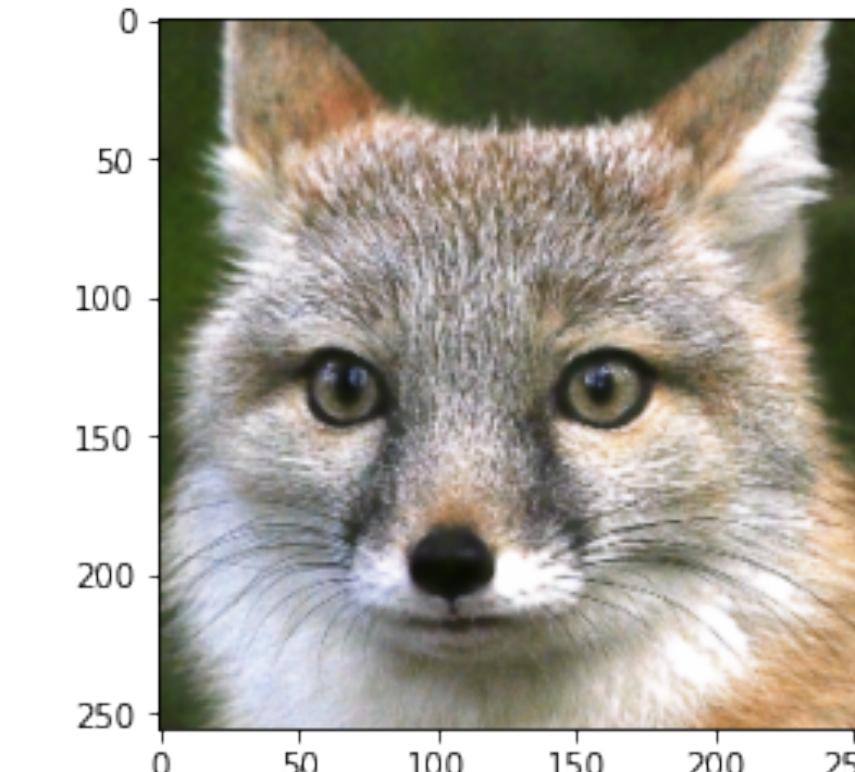
**input RGB**



**basic coordinate  
ReLU-MLP**



**best effort coordinate  
ReLU-MLP**  
(pre-trained with  
periodic embedding)



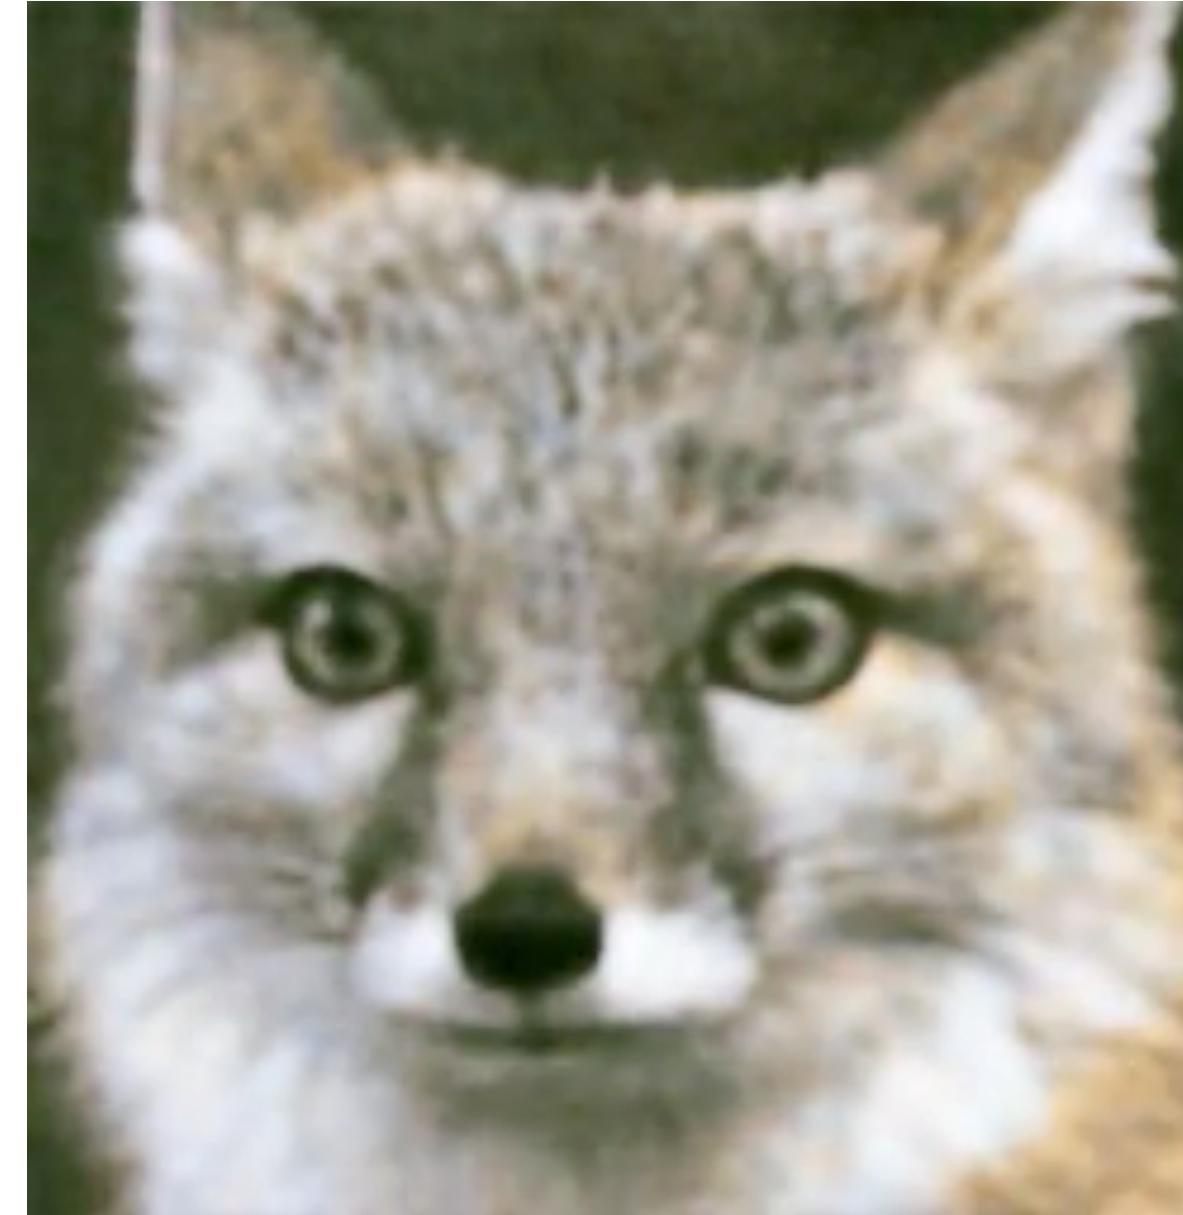
**(Gaussian) Fourier  
feature mapping**

remaining issue, the dimensions and range (std.dev.) of the basis matrix  $\mathbf{B}$  needs to be carefully chosen

$$\gamma(\mathbf{v}) = [\cos(2\pi \mathbf{B}\mathbf{v}), \sin(2\pi \mathbf{B}\mathbf{v})]^T$$

each entry in  $\mathbf{B} \in \mathbb{R}^{m \times d}$  is sampled from  $\mathcal{N}(0, \sigma^2)$

# Evolution of iterations for different input features



**exp-Fourier #net2 263'683**



**#hash 70'770 #net 206'595**

→ hash-grid encoding not covered today due to time constraints (feel free to ask)  
Thomas Müller et al.: **Instant Neural Graphics Primitives with a Multiresolution Hash Encoding**  
ACM Transactions on Graphics 2022 <https://nvlabs.github.io/instant-ngp>

# Representing images as functions with periodic activations

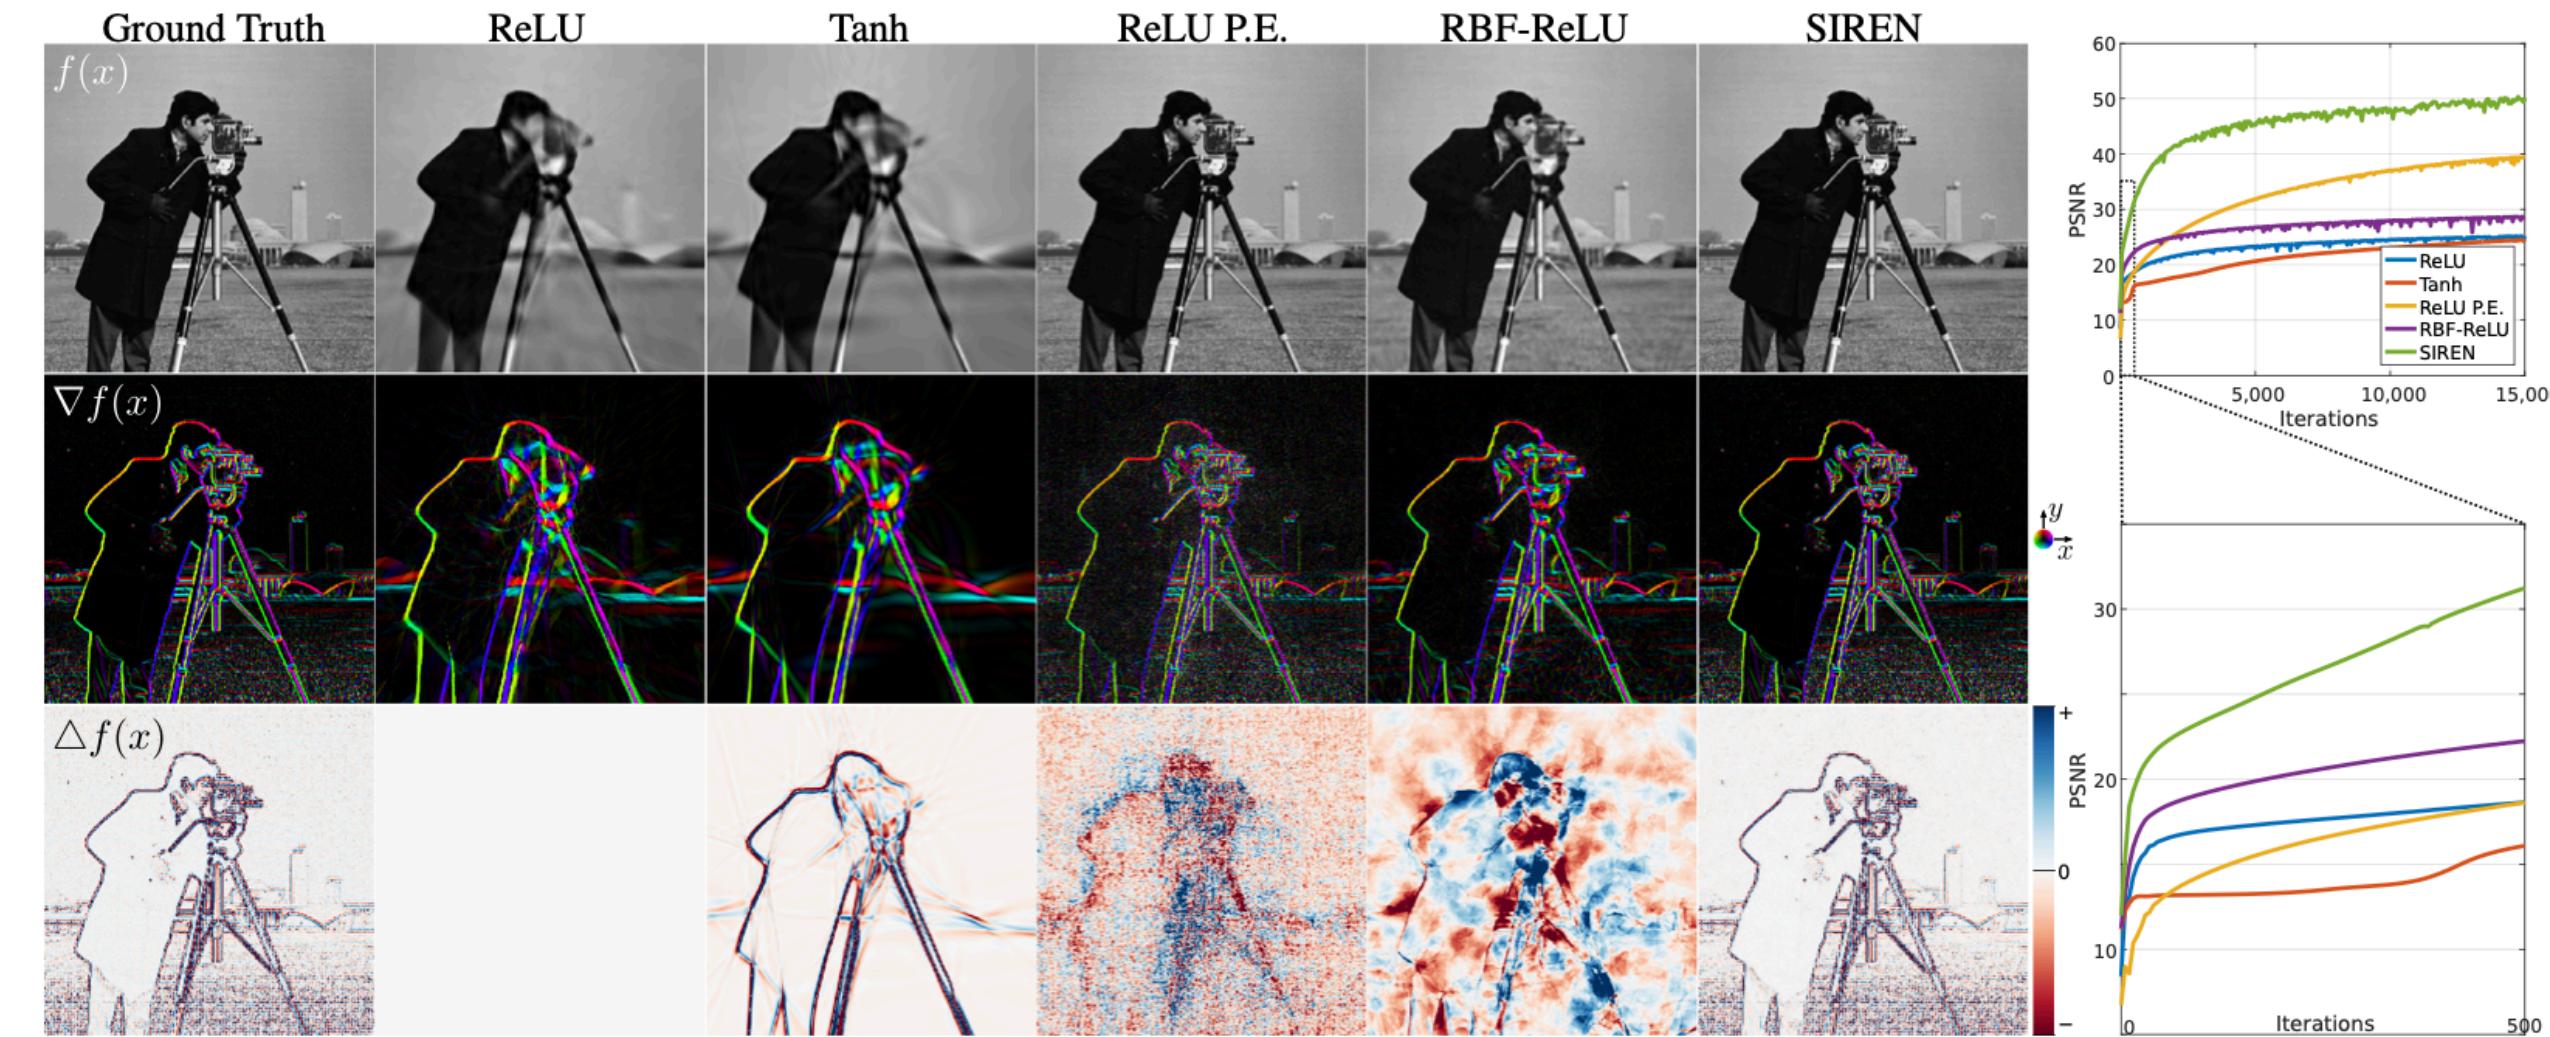


Figure 1: Comparison of different neural network architectures fitting the implicit representation of an image (ground truth: top left). The representation is only supervised on the target image but we also show first- and second-order derivatives of the function fit in rows 2 and 3, respectively.

Vincent Sitzmann\*, Julien N. P. Martel\* et al.: "**Implicit Neural Representations with Periodic Activation Functions**"

NeurIPS 2020

<https://vsitzmann.github.io/siren/>

→ **introduces sinusoidal activations**  
networks are usually called SIRENs

# Theoretical formulation

We are interested in a class of functions  $\Phi$  that satisfy equations of the form

$$F(\mathbf{x}, \Phi, \nabla_{\mathbf{x}}\Phi, \nabla_{\mathbf{x}}^2\Phi, \dots) = 0, \quad \Phi : \mathbf{x} \mapsto \Phi(\mathbf{x}). \quad (1)$$

This implicit problem formulation takes as input the spatial or spatio-temporal coordinates  $\mathbf{x} \in \mathbb{R}^m$  and, optionally, derivatives of  $\Phi$  with respect to these coordinates. Our goal is then to learn a neural network that parameterizes  $\Phi$  to map  $\mathbf{x}$  to some quantity of interest while satisfying the constraint presented in Equation (1). Thus,  $\Phi$  is implicitly defined by the relation defined by  $F$  and we refer to neural networks that parameterize such implicitly defined functions as *implicit neural representations*.

Our goal is to solve problems of the form presented in Equation (1). We cast this as a feasibility problem, where a function  $\Phi$  is sought that satisfies a set of  $M$  constraints  $\{\mathcal{C}_m(\mathbf{a}(\mathbf{x}), \Phi(\mathbf{x}), \nabla\Phi(\mathbf{x}), \dots)\}_{m=1}^M$ , each of which relate the function  $\Phi$  and/or its derivatives to quantities  $\mathbf{a}(\mathbf{x})$ :

$$\text{find } \Phi(\mathbf{x}) \text{ subject to } \mathcal{C}_m(\mathbf{a}(\mathbf{x}), \Phi(\mathbf{x}), \nabla\Phi(\mathbf{x}), \dots) = 0, \forall \mathbf{x} \in \Omega_m, m = 1, \dots, M \quad (2)$$

This problem can be cast in a loss function that penalizes deviations from each of the constraints on their domain  $\Omega_m$ :

$$\mathcal{L} = \int_{\Omega} \sum_{m=1}^M \mathbf{1}_{\Omega_m}(\mathbf{x}) \|\mathcal{C}_m(\mathbf{a}(\mathbf{x}), \Phi(\mathbf{x}), \nabla\Phi(\mathbf{x}), \dots)\| d\mathbf{x}, \quad (3)$$

with the indicator function  $\mathbf{1}_{\Omega_m}(\mathbf{x}) = 1$  when  $\mathbf{x} \in \Omega_m$  and 0 when  $\mathbf{x} \notin \Omega_m$ . In practice, the loss function is enforced by sampling  $\Omega$ . A dataset  $\mathcal{D} = \{(\mathbf{x}_i, \mathbf{a}_i(\mathbf{x}))\}_i$  is a set of tuples of coordinates  $\mathbf{x}_i \in \Omega$  along with samples from the quantities  $\mathbf{a}(\mathbf{x}_i)$  that appear in the constraints. Thus, the loss in Equation (3) is enforced on coordinates  $\mathbf{x}_i$  sampled from the dataset, yielding the loss  $\tilde{\mathcal{L}} = \sum_{i \in \mathcal{D}} \sum_{m=1}^M \|\mathcal{C}_m(a(\mathbf{x}_i), \Phi(\mathbf{x}_i), \nabla\Phi(\mathbf{x}_i), \dots)\|$ . In practice, the dataset  $\mathcal{D}$  is sampled dynamically at training time, approximating  $\mathcal{L}$  better as the number of samples grows, as in Monte Carlo integration.

We parameterize functions  $\Phi_\theta$  as fully connected neural networks with parameters  $\theta$ , and solve the resulting optimization problem using gradient descent.

**continuous implicit neural representation** using periodic activation functions **fit complicated signals**, such as natural images and 3D shapes, **and their derivatives robustly**.

We propose SIREN, a simple neural network architecture for implicit neural representations that uses the sine as a periodic activation function:

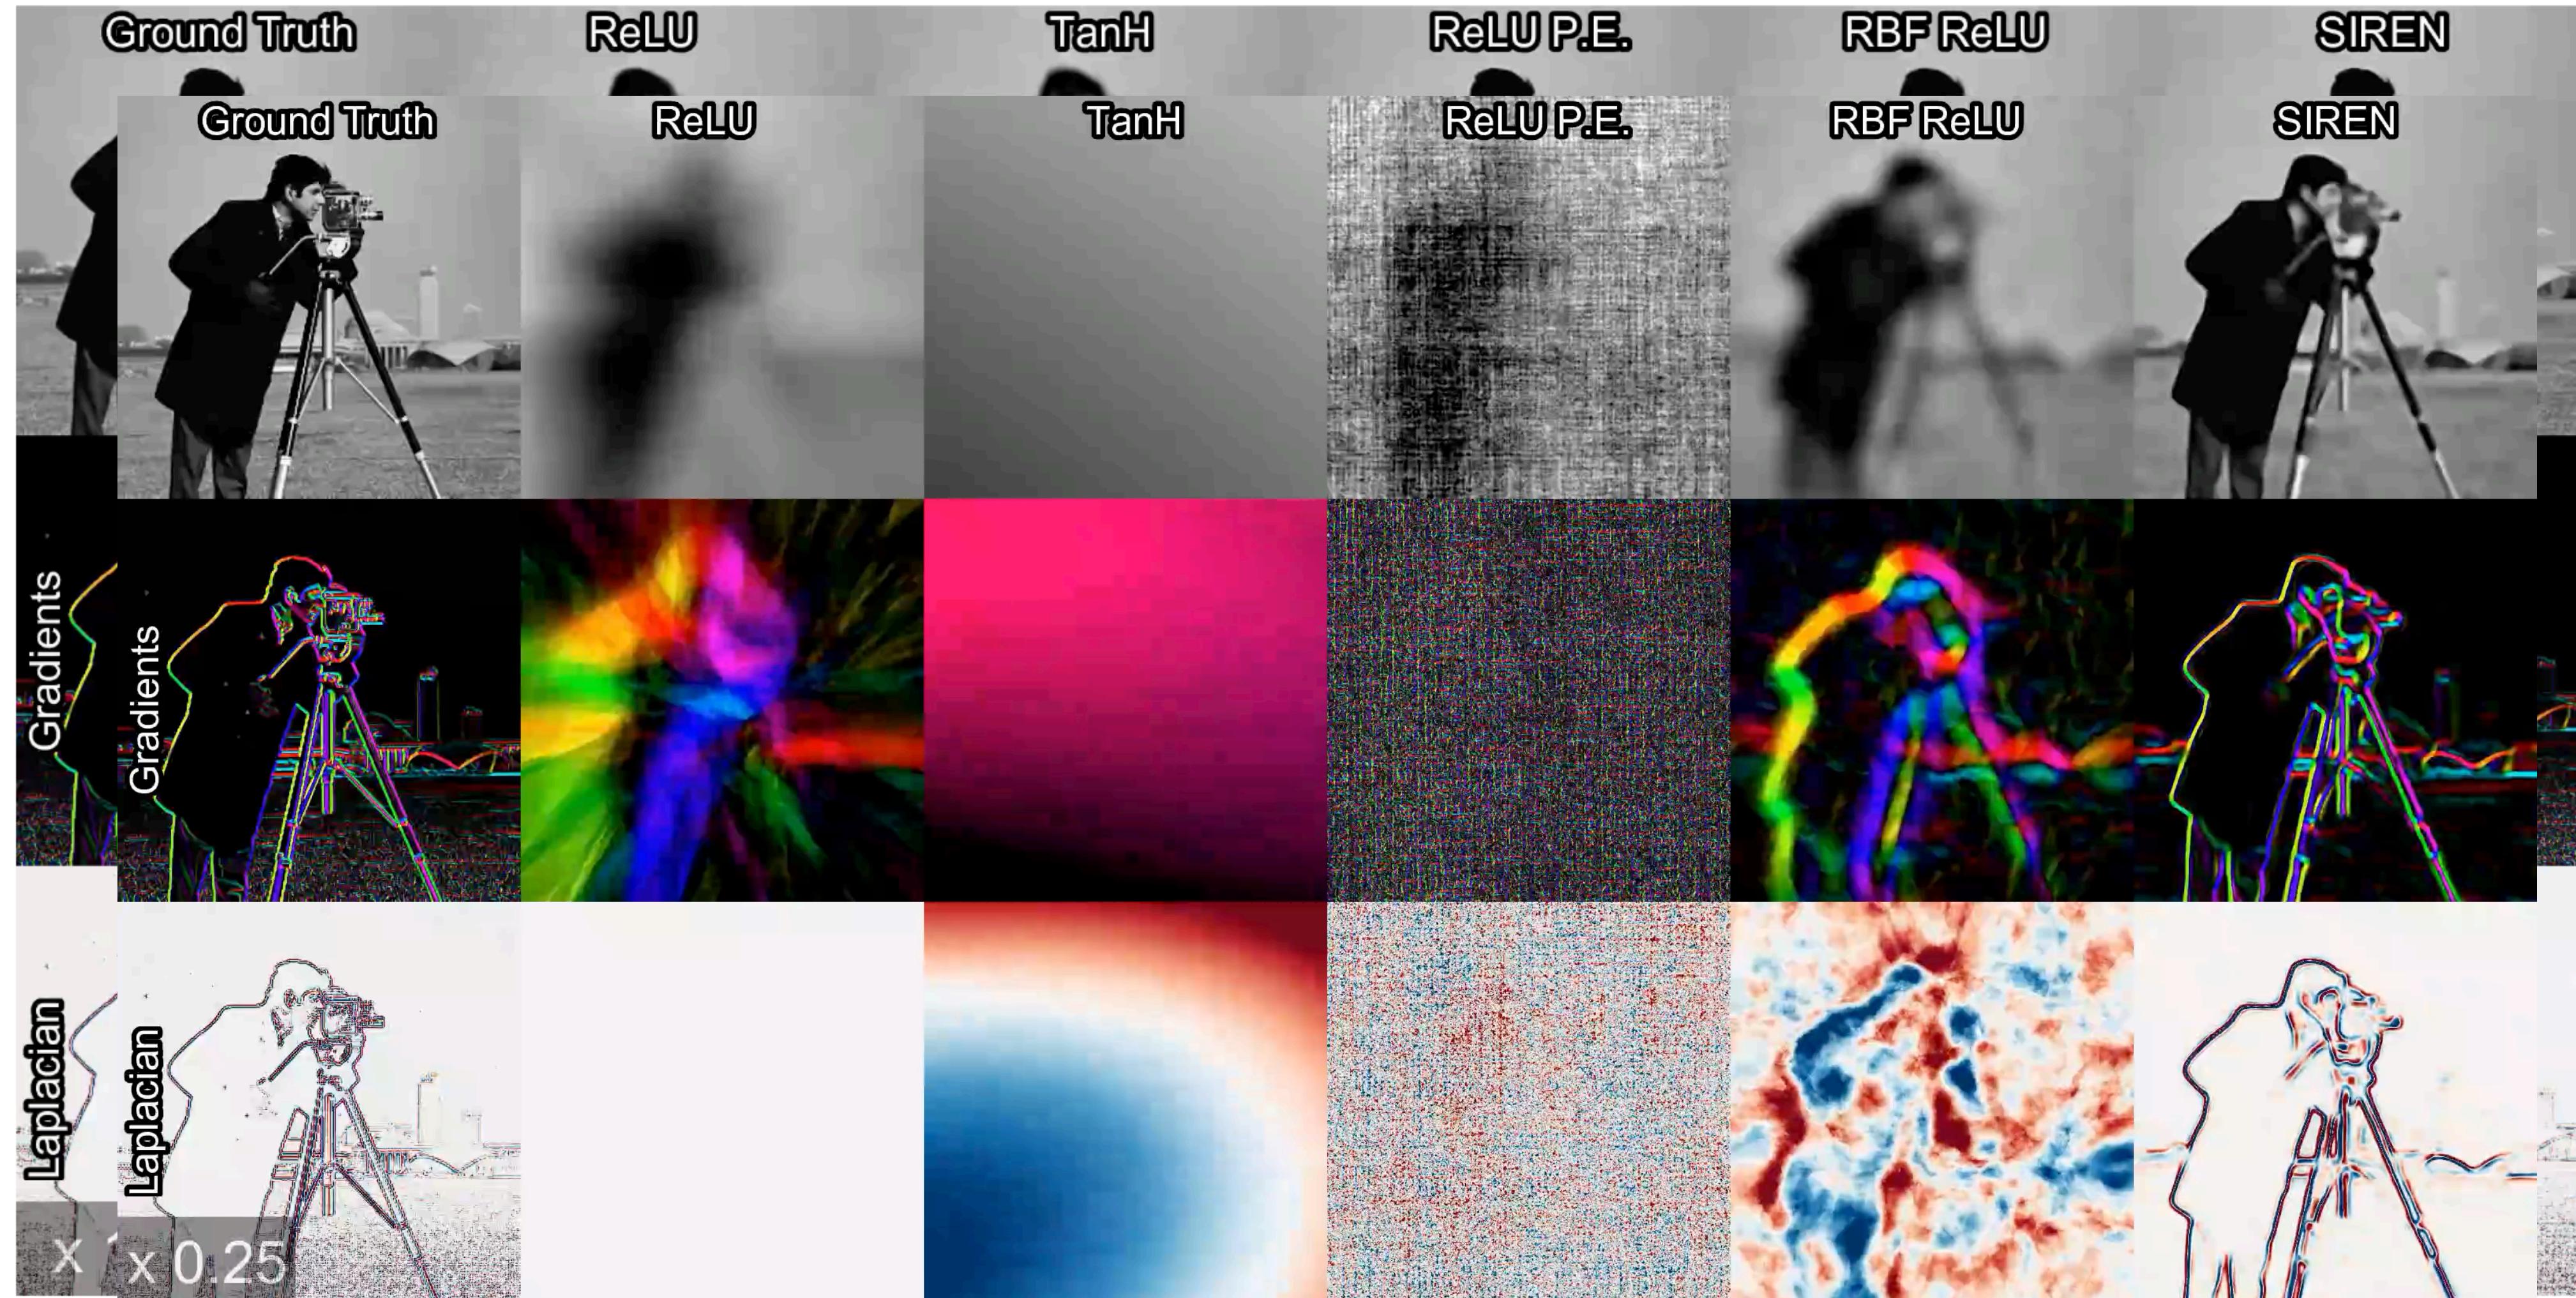
$$\Phi(\mathbf{x}) = \mathbf{W}_n (\phi_{n-1} \circ \phi_{n-2} \circ \dots \circ \phi_0)(\mathbf{x}) + \mathbf{b}_n, \quad \mathbf{x}_i \mapsto \phi_i(\mathbf{x}_i) = \sin(\mathbf{W}_i \mathbf{x}_i + \mathbf{b}_i). \quad (4)$$

Here,  $\phi_i : \mathbb{R}^{M_i} \mapsto \mathbb{R}^{N_i}$  is the  $i^{th}$  layer of the network. It consists of the affine transform defined by the weight matrix  $\mathbf{W}_i \in \mathbb{R}^{N_i \times M_i}$  and the biases  $\mathbf{b}_i \in \mathbb{R}^{N_i}$  applied on the input  $\mathbf{x}_i \in \mathbb{R}^{M_i}$ , followed by the sine nonlinearity applied to each component of the resulting vector.

Interestingly, any derivative of a SIREN *is itself a SIREN*, as the derivative of the sine is a cosine, i.e., a phase-shifted sine (see supplemental). Therefore, the derivatives of a SIREN inherit the properties of SIRENs, enabling us to supervise any derivative of SIREN with “complicated” signals. In our experiments, we demonstrate that when a SIREN is supervised using a constraint  $\mathcal{C}_m$  involving the derivatives of  $\phi$ , the function  $\phi$  remains well behaved, which is crucial in solving many problems, including boundary value problems (BVPs).

→ **some caveats: initialisation can be tricky, function has to be fitted at test-time**

# Implicit Neural Representations with Periodic Activation Functions

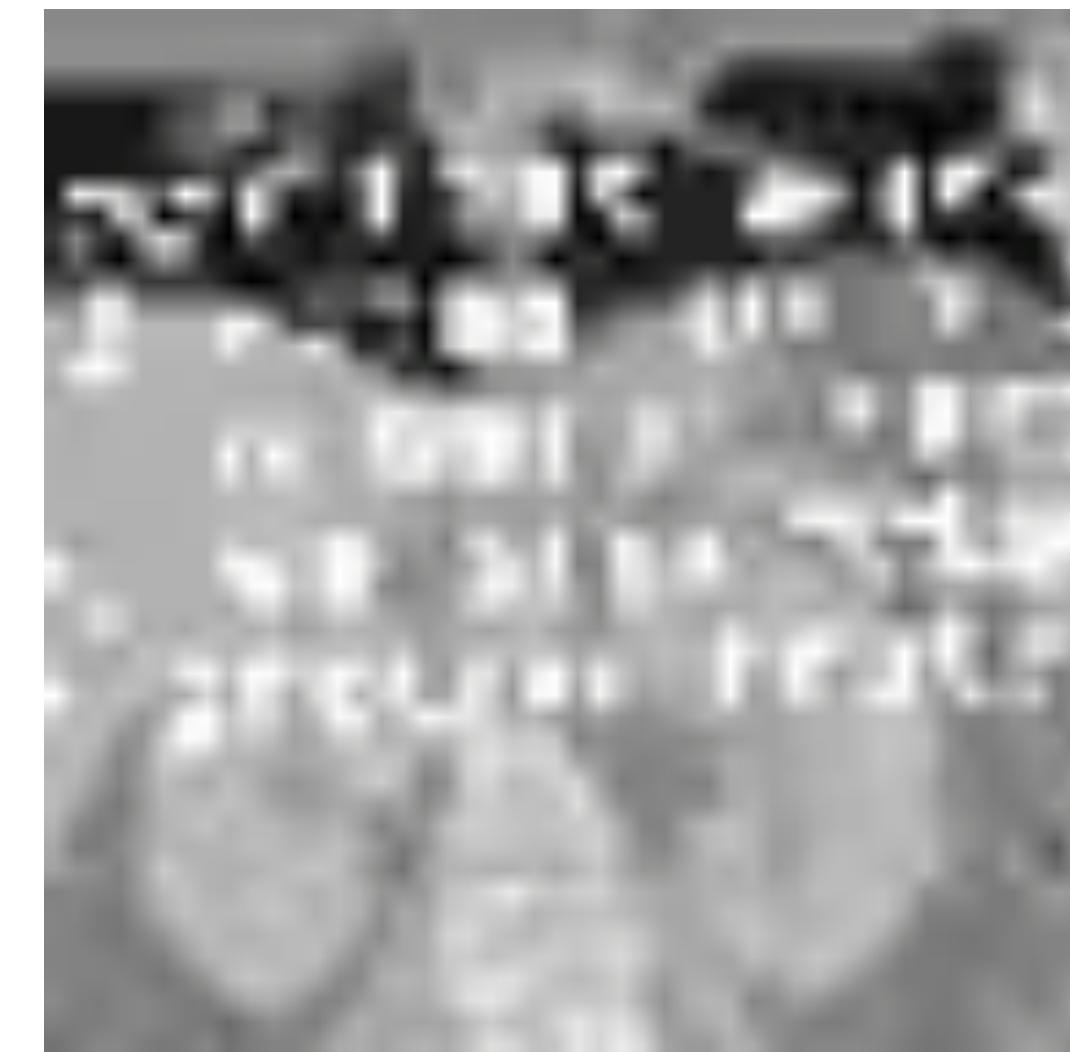


# Representing images as functions with periodic activations

**typical forward pass for MLP**, first layer has to have sinusoidal (or exponential) activation, the other ones could be ReLUs

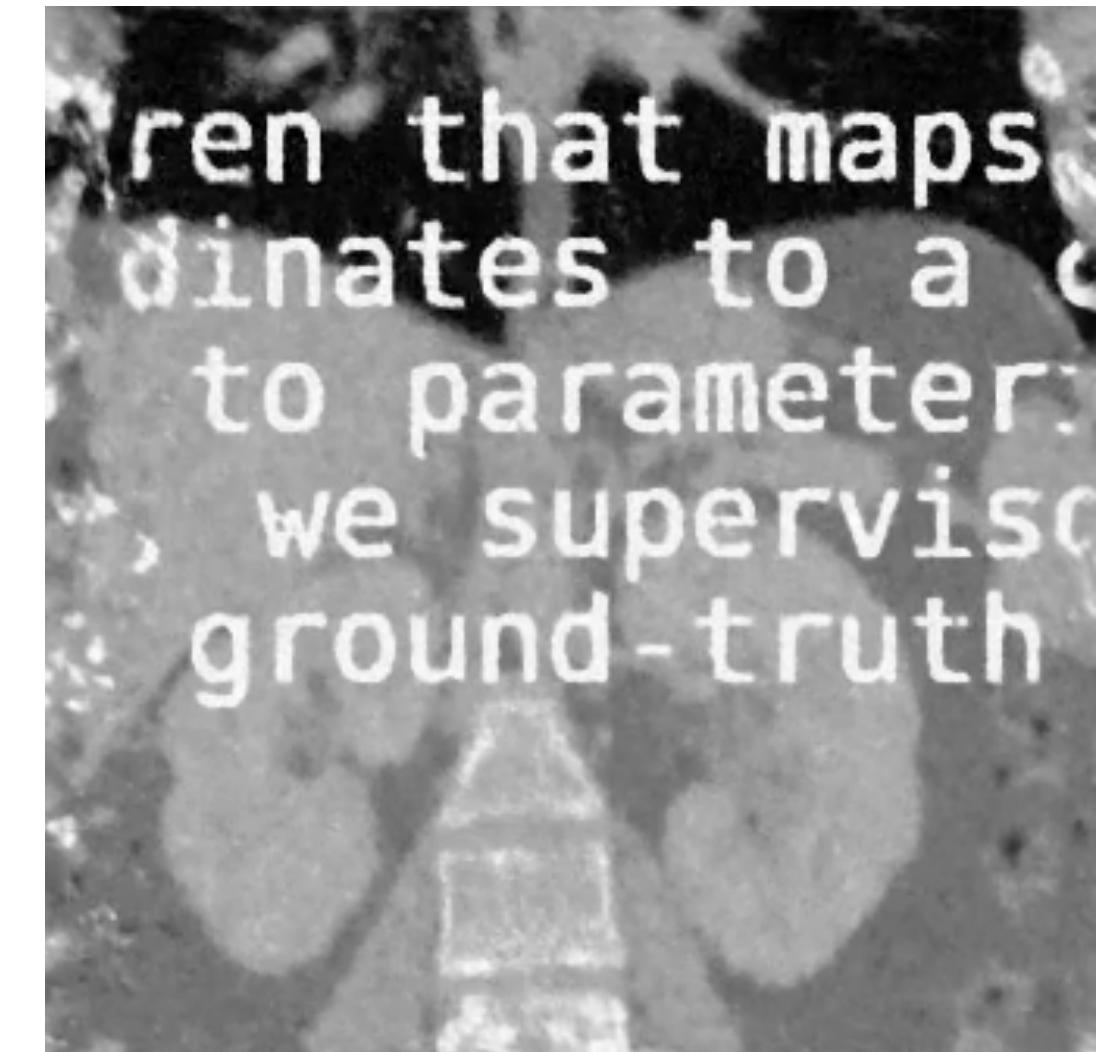
```
def forward(self,x):
    scale = self.scale
    x = torch.sin(scale*self.net[0](x))
    for i in range(1,len(self.net)-1):
        #x = torch.exp(-scale*self.net[i](x).pow(2))
        #x = F.relu(self.net[i](x))
        x = torch.sin(scale*self.net[i](x))
    return self.net[-1](x)
```

→ SIREN paper discusses **multiple advanced applications**, including inpainting, solving wave equations, and representing shapes as signed distance function



**input sequence is low resolution (84x84 pixels)**  
"camera position" of 64 frames is known

same example from beginning



SIREN coordinate-MLP **reconstructs 8x super-resolved image** with dimensions 334x334

# Hands-on Tutorial

[ziadhemidi.github.io/INR4BVM25.github.io](https://ziadhemidi.github.io/INR4BVM25.github.io)